

Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA)
Universitat Politècnica de València

Part 4: Memory management

Unit 12

Virtual memory (II)

fSO

DISCA



UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA

- **Goals**

- To understand **2nd chance** page replacement algorithm as an LRU approximation
- To know the **thrashing** problem and its solutions
- To analyse memory **frame management** techniques

- **Bibliography**

- Silberschatz, chapter 9

- **LRU review: problem**
- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- Frame reservation



- **Optimal replacement algorithm**
 - Victim page: the one which take longer to be referenced
 - Minimum number of faults -> impossible to implement (future is unknown)
- **LRU (Least Recently Used).** [Optimal approximation]
 - Victim page: the one that lasted more without being referenced
 - It is a stack algorithm
 - Implementation
 - Using counters: Every page has associated a counter or timestamp
 - Using an ordered list of referenced pages: When a page is referenced it is moved to the end. The victim page is the first one
 - Using a stack: when a page is referenced, it is added on top. The victim page is the last one. As a referenced page could be in the middle of the stack, it is better to use a double linked list with a head and tail pointer.

Exercise 4. Replacement scope

On a virtual memory system, with 1024 byte page size, the OS has allocated 6 frames (from 0 to 5) to two processes A y B. At time $t = 10$, A and B page tables have the following content:

Process A page table	Frame	Valid bit	Counter
	0	i	
	1	i	
	2	v	10
	3	v	3
	4	i	
	5	v	5
	6	i	
	7	i	

Process B page table	Frame	Valid bit	Counter
	0	i	
	1	i	
	2	i	
	3	v	2
	4	i	
	5	i	
	6	i	
	7	i	

Then the processes emit the following logical address sequence. Consider that all the addresses are legal:

A,100; A,4000; B,100; A,7000; B,2100; B,1028; A,5800; A,100

Obtain what pages are allocated on every frame and the physical address translation of the first and the last access in the following situations:

- The replacement algorithm is **LRU** with **global scope**
- The replacement algorithm is **LRU** with **local scope**. Process A has 4 frames and process B has 2 frames

Exercise 4. Replacement scope

On a virtual memory system, with 1024 byte page size, the OS has allocated 6 frames (from 0 to 5) to two processes A y B. At time $t = 10$, A and B page tables have the following content:

A and B have 8 pages
1 page 1024 bytes = 2^{10}
Size $8 * 1024 = 8192$
Number of bits: $2^{13} = 8192$
Logical address: 3 + 10 bits

Process A page table

	Frame	Valid bit	Counter
0		i	
1		i	
2	2	v	10
3	5	v	3
4		i	
5	4	v	5
6		i	
7		i	

Process B page table

	Frame	Valid bit	Counter
0		i	
1		i	
2		i	
3	1	v	2
4		i	
5		i	
6		i	
7		i	

a) The replacement algorithm is **LRU** with **global scope**

	A	A	B	A	B	B	A	A
L.Add	100	4000	100	7000	2100	1028	5800	100
Page	0	3	0	6	2	1	5	0
offset	100	928	100	856	52	4	680	100

frame	t = 10	11	12	13	14	15	16	17	18
0	free	A0, 11	A0, 11	A0, 11	A0, 11	A0, 11	A0, 11	A5, 17	A5, 17
1	B3, 2	B3, 2	B3, 2	B3, 2	A6, 14	A6, 14	A6, 14	A6, 14	A6, 14
2	A2, 10	A2, 10	A2, 10	A2, 10	A2, 10	A2, 10	B1, 16	B1, 16	B1, 16
3	free	free	free	B0, 13	B0, 13	B0, 13	B0, 13	B0, 13	B0, 13
4	A5, 5	A5, 5	A5, 5	A5, 5	A5, 5	B2, 15	B2, 15	B2, 15	B2, 15
5	A3, 3	A3, 3	A3, 12	A3, 12	A3, 12	A3, 12	A3, 12	A3, 12	A0, 18

Exercise 4. Replacement scope

On a virtual memory system, with 1024 byte page size, the OS has allocated 6 frames (from 0 to 5) to two processes A y B. At time $t = 10$, A and B page tables have the following content:

A and B have 8 pages
1 page 1024 bytes = 2^{10}
Size $8 * 1024 = 8192$
Number of bits: $2^{13} = 8192$
Logical address: 3 + 10 bits

Process A page table

	Frame	Valid bit	Counter
0		i	
1		i	
2	2	v	10
3	5	v	3
4		i	
5	4	v	5
6		i	
7		i	

Process B page table

	Frame	Valid bit	Counter
0		i	
1		i	
2		i	
3	1	v	2
4		i	
5		i	
6		i	
7		i	

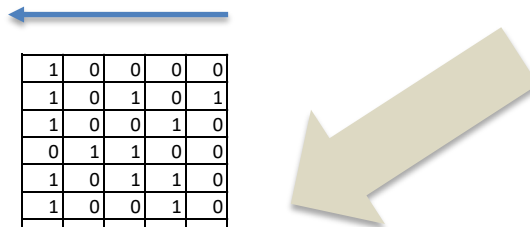
a) The replacement algorithm is **LRU** with **global local**

	A	A	B	A	B	B	A	A
L.Add	100	4000	100	7000	2100	1028	5800	100
Page	0	3	0	6	2	1	5	0
offset	100	928	100	856	52	4	680	100

frame	t = 10	11	12	13	14	15	16	17	18
0	free	A0, 11	A0, 11	A0, 11	A0, 11	A0, 11	A0, 11	A0, 11	A0, 18
1	B3, 2	B3, 2	B3, 2	B3, 2	B3, 2	B2, 15	B2, 15	B2, 15	B2, 15
2	A2, 10	A2, 10	A2, 10	A2, 10	A2, 10	A2, 10	A2, 10	A5, 17	A5, 17
3	free	free	free	B0, 13	B0, 13	B0, 13	B1, 16	B1, 16	B1, 16
4	A5, 5	A5, 5	A5, 5	A5, 5	A6, 14	A6, 14	A6, 14	A6, 14	A6, 14
5	A3, 3	A3, 3	A3, 12	A3, 12	A3, 12	A3, 12	A3, 12	A3, 12	A3, 12


- LRU review: problem
- **2nd chance replacement algorithm**
- Frame allocation
- Thrashing
- Frame reservation

- Supporting LRU as replacement algorithm is too expensive
- Solution:** To approximate LRU using the reference bit
- Reference bit** can be set by the hardware every time a page is referenced.
- Scheme:** Every 100ms, the OS copies the reference bits and set to 0.



1	0	0	0	0
1	0	1	0	1
1	0	0	1	0
0	1	1	0	0
1	0	1	1	0
1	0	0	1	0
1	1	0	1	1
1	0	1	0	1
1	0	1	0	1
0	1	1	0	0
1	0	0	1	1

victim




- Supporting LRU as replacement algorithm is too expensive
- Solution:** To approximate LRU using the reference bit
- Reference bit can be set by the hardware every time a page is referenced.

2nd chance algorithm

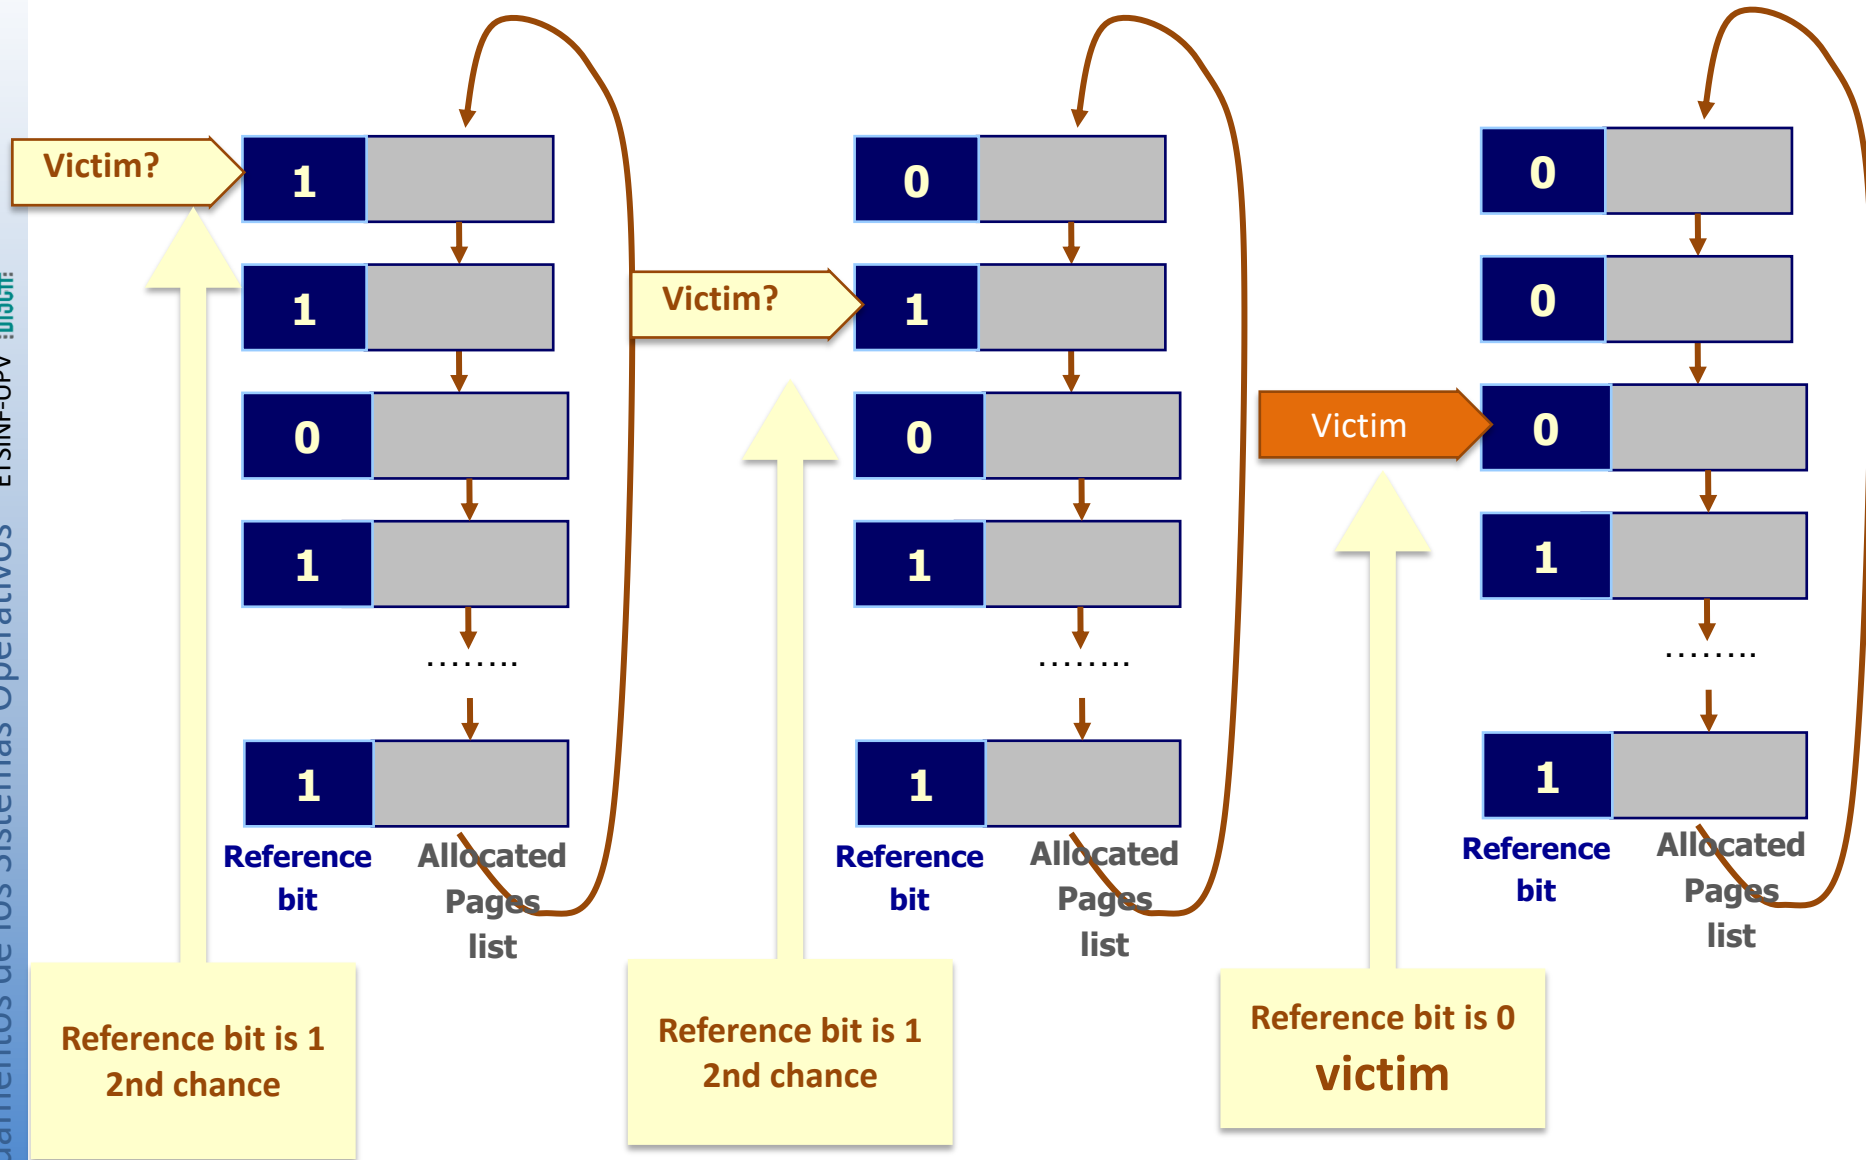
- It reviews the page reference bit on FIFO ordering
- If the page reference bit is 1**
 - Put reference bit to 0
 - Let page in memory
- If the page reference bit is 0**
 - The page is chosen as victim
 - Next page is the first candidate for next victim

Next victim?

2nd
opportunity




2nd chance replacement algorithm



Example: Main memory with 4 frames

T	0	1	2	3	4	5	6	7	8	9	10
Reference string	1	2	3	4	1	2	5	2	3	1	2
Frame 0	1	1	1	1	1						
Frame 1		2	2	2	2						
Frame 2			3	3	3						
Frame 3				4	4						


	1	1	1	1	1	1	1	1	1
	0		1	2	1	2	1	2	1
	0		0		1	3	1	3	1
	0		0		1	4	1	4	1
Time	0	1	2	3	4				



Example: Main memory with 4 frames

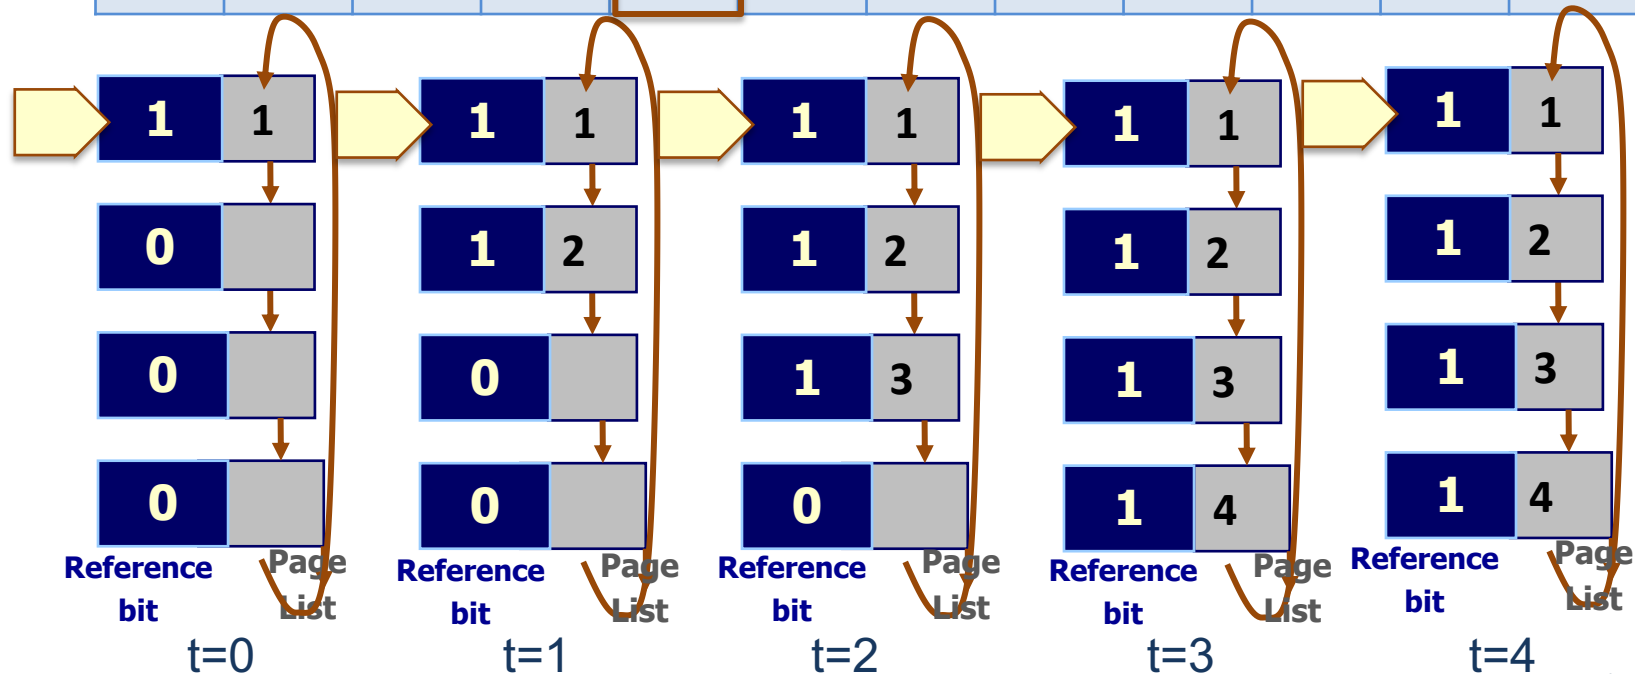
T	0	1	2	3	4	5	6	7	8	9	10
Reference string	1	2	3	4	1	2	5	2	3	1	2
Frame 0	1	1	1	1	1	1	5	5	5	5	5
Frame 1		2	2	2	2	2	2	2	2	2	2
Frame 2			3	3	3	3	3	3	3	3	3
Frame 3				4	4	4	4	4	4	1	1

	1	1	1	1	1	1	1	1	1	
	0		1	2	1	2	1	2	1	2
	0		0		1	3	1	3	1	3
	0		0		1	4	1	4	1	4
Time	0	1	2	3	4					



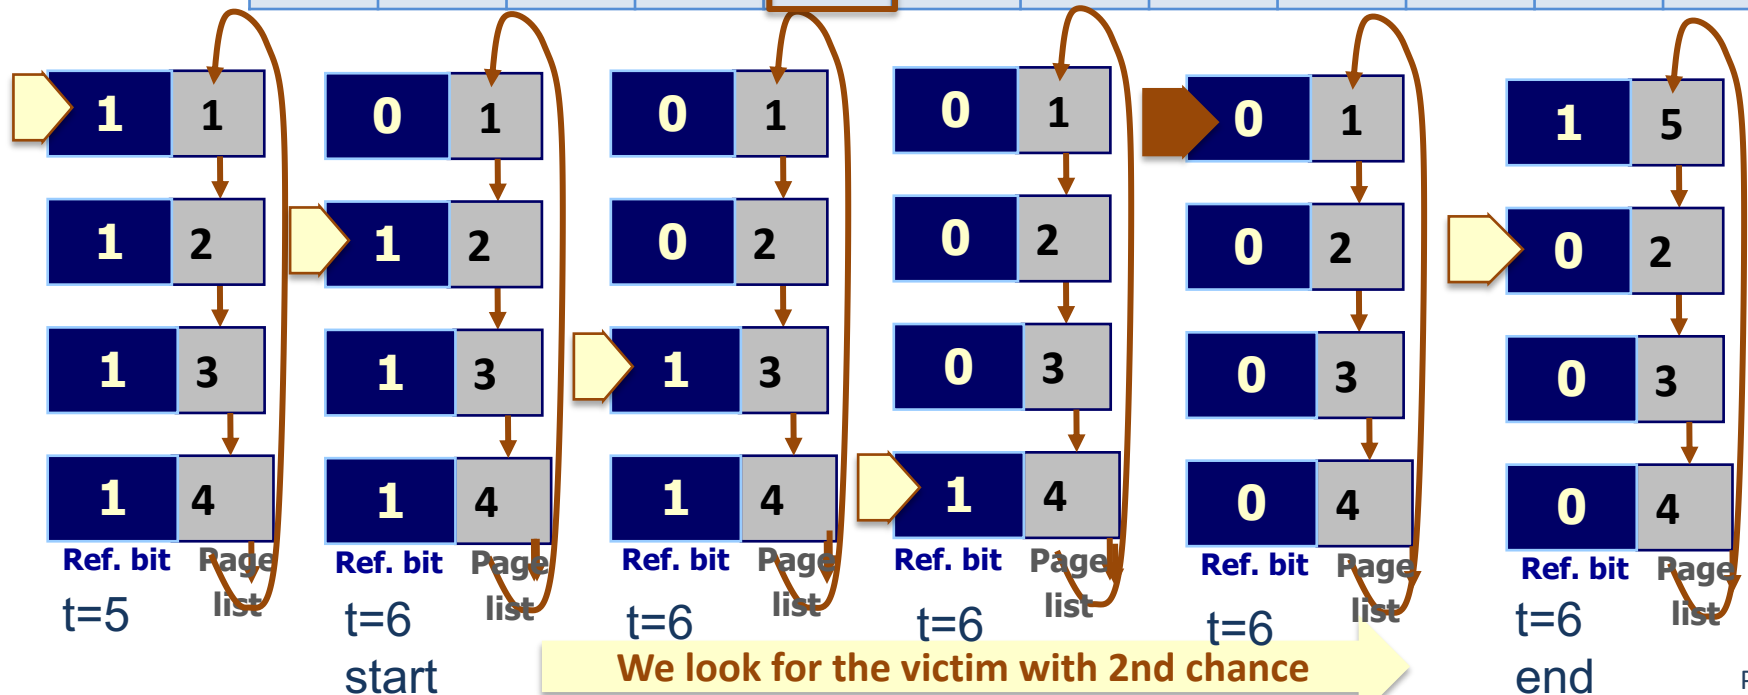
Example: Main memory with 4 frames

T	0	1	2	3	4	5	6	7	8	9	10
Reference string	1	2	3	4	1	2	5	2	3	1	2
Frame 0	1	1	1	1	1						
Frame 1		2	2	2	2						
Frame 2			3	3	3						
Frame 3				4	4						



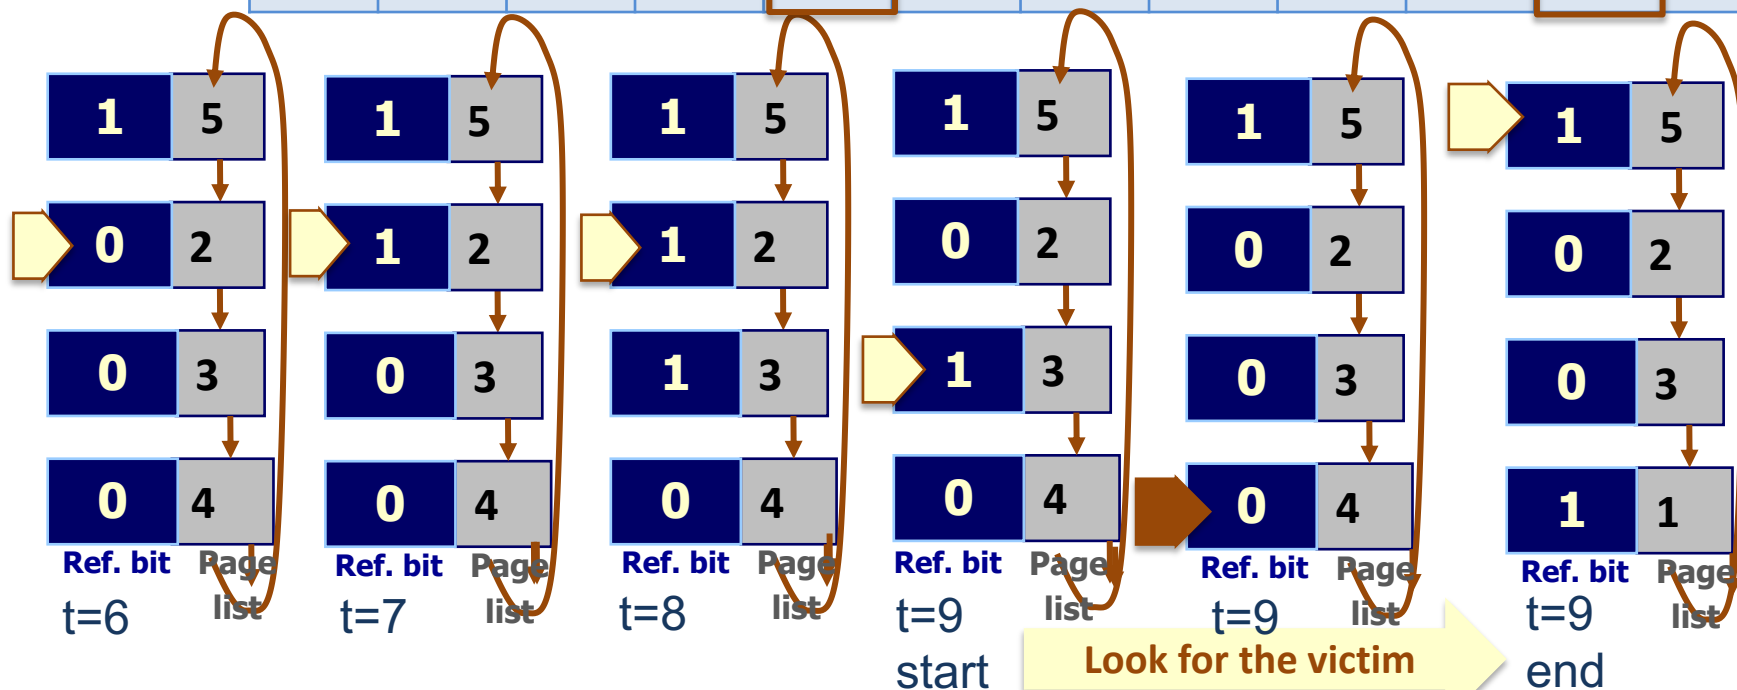
Example: Main memory with 4 frames

T	0	1	2	3	4	5	6	7	8	9	10
Reference string	1	2	3	4	1	2	5	2	3	1	2
Frame 0	1	1	1	1	1	1	5				
Frame 1		2	2	2	2	2	2				
Frame 2			3	3	3	3	3				
Frame 3				4	4	4	4				



Example: Main memory with 4 frames

T	0	1	2	3	4	5	6	7	8	9	10
Reference string	1	2	3	4	1	2	5	2	3	1	2
Frame 0	1	1	1	1	1	1	5	5	5	5	
Frame 1		2	2	2	2	2	2	2	2	2	
Frame 2			3	3	3	3	3	3	3	3	
Frame 3				4	4	4	4	4	4	1	



There exist several algorithms

- **The least frequently used (LFU)** page-replacement algorithm requires that the page with **the smallest count be replaced**. The reason for this selection is that an actively used page should have a large reference count. A problem arises, however, when a page is used heavily during the initial phase of a process but then is never used again.
- **The most frequently used (MFU)** page-replacement algorithm is based on the argument that the page with the smallest count was probably just brought in and has yet to be used.

- 2nd chance replacement algorithm
- **Frame allocation**
- Thrashing
- Frame reservation

Question:

- How do we allocate the fixed amount of free memory among the various processes?
- If we have 38 free frames and 2 processes, how many frames does each process get?
- Simple case: 128 frames
 - OS: 35
 - User processes: 93

- Frame allocation problem
 - Free frame list:
 - Frame management requires a **data structure** where **free frames** are kept
 - Frame to process delivery policy and the OS
 - The OS gets the required number of frames to execute itself
 - **Processes receive the minimum initial number of frames** and the remaining ones on demand
 - The minimum number of initially assigned frames depends on the indirection level in the CPU instruction set → To execute an instruction all its operands must be allocated in main memory

What the hell is this?

- Consider a machine in which all memory-reference instructions may reference two memory addresses.
- In this case, we need at least **one frame for the instruction** and **two frames for the memory references**.
 - `ADD R1, A, B R1 <- M[A] + M[B]`
- Additionally, if **one-level indirect addressing is allowed** (for example, a load instruction on frame 16 can refer to an address on other frame, which is an indirect reference to another), then **paging requires at least five frames** per process.
- Depending on the instruction size, an instruction could be allocated in the limit of the page and the next one
- **6 pages could be required** (minimum number of pages).

- Frame distribution policies:

(Considering **m frames** and **n processes**)

- **Fair allocation:** All processes allocate A_i frames equally

$$A_i = \left\lfloor \frac{m}{n} \right\rfloor \quad \text{Remainder to a frame free pool}$$

- **Proportional allocation:** A process P_i with size S_i allocate A_i frames computed as:


$$A_i = \frac{S_i}{\sum S_i} * m$$

- **Priority allocation:** higher priority processes allocate more frames. If a high priority process generates a page fault, the victim can be one of its frames or a frame of a process with lower priority.

- **Replacement policy scope**
 - **Local replacement:** only pages allocated to the process that generates the page fault can be replaced
 - It cannot choose a victim from another process
 - The number of process allocated frames does not change
 - A process execution is not affected by the remaining processes
 - Advantage: Sensible response time
 - Disadvantage: Worse global memory management
 - **Global replacement:** the victim is chosen between all allocated pages in main memory
 - The victim can belong to another process
 - The number of process allocated frames can change
 - Disadvantage: Response time sensitive to system load
 - Advantage: Better global memory management

- 2nd chance replacement algorithm
- Frame allocation
- **Thrashing**
- Frame reservation

- The thrashing problem:
 - Memory becomes scarce and processes generate **a lot of page faults**
→ I/O time becomes dominant
 - The OS allows more processes entering for execution regarding the **low level of CPU use**
 - This makes things worse because the same amount of memory has to be shared with more and more processes → page fault rate keeps increasing




How to solve thrashing?

- To anticipate and to prevent the problem
 - Using a working set model
 - Controlling page fault rate
- Once thrashing is detected
 - Swap out processes with a medium term scheduler

A process is thrashing if it is spending more time paging than executing => performance problems

- Reference locality principle
 - Instructions and data processed recently (and the ones close to them) have a high probability of being processed in the near future
 - Locality:
 - Set of pages that a process uses as a whole
 - It is hard to identify
 - Thrashing happens when
locality sizes > total main memory size




- **Working set model:** preventive technique
 - It assumes **reference locality principle**
 - Obtain the **number of pages that a process needs to have simultaneously in memory** to avoid thrashing
 - **Working set (WS):**
 - Set of pages accessed in a process: logical addresses last referenced
 - **Working set window (WS):**
 - Fixed consecutive number of references Δ used to compute WS
 - WS is built with the set of pages accessed in the last Δ references
 - **Working set size (WSS):** Number of different pages that belong to the WS. $WSS_i \Rightarrow$ number of pages of the WS of the process P_i
 - In a system with **m frames and n processes** $P_1 \dots P_n$ there is thrashing when **$\sum WSS_i > m$**

- **Example of WS model**

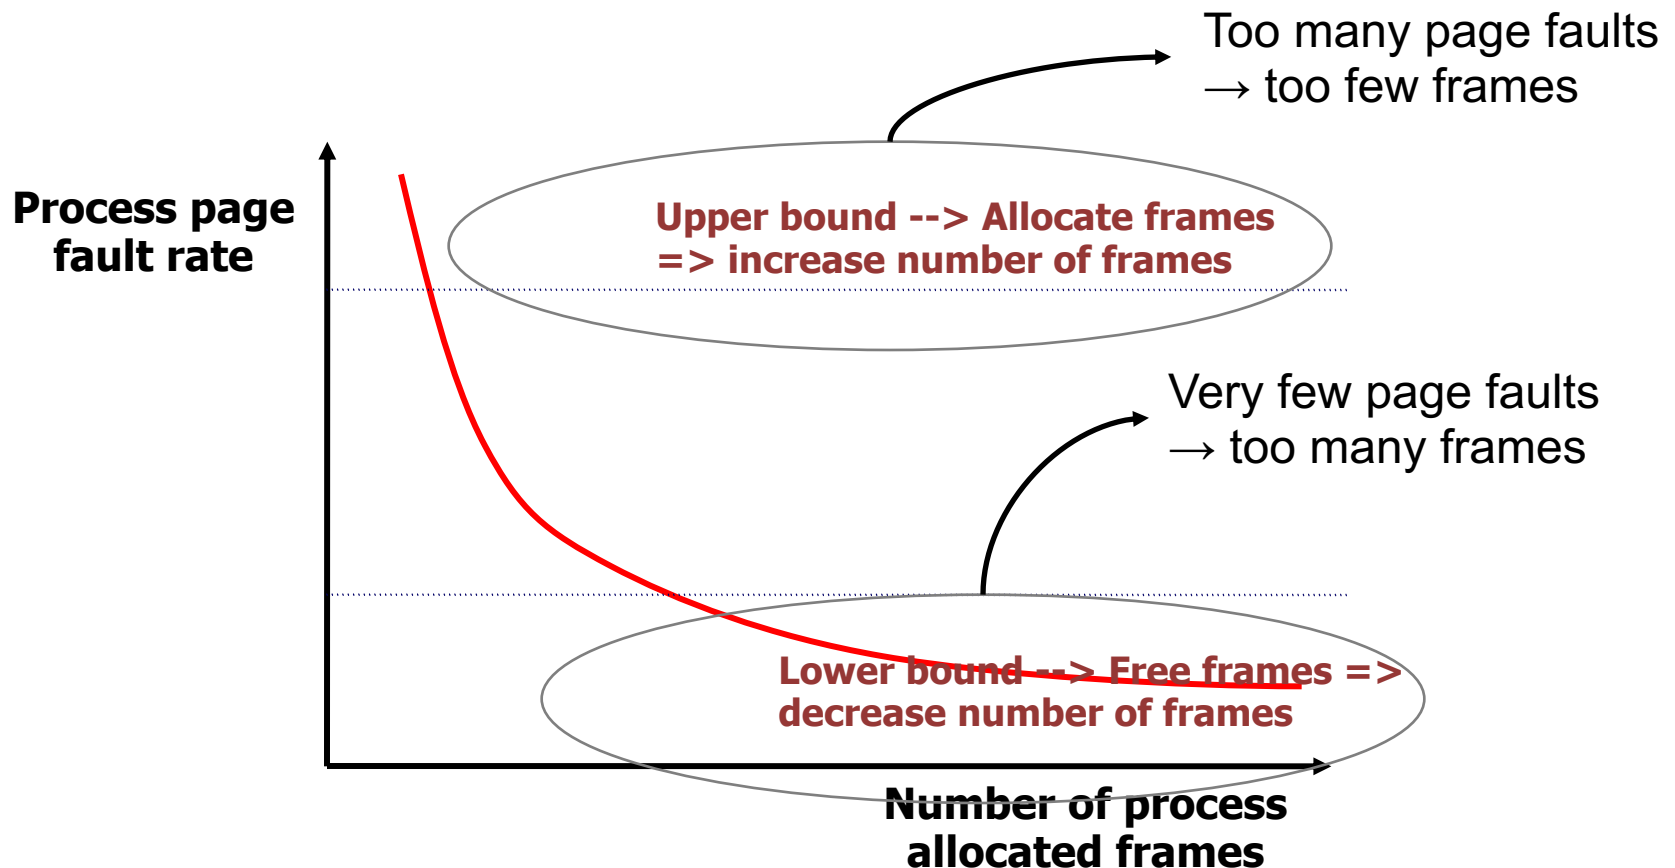
- WS window $\Delta = 9$
- Compute WS and WSS in $t1$ and $t2$

Sequence of referenced pages



- **Page fault rate control**

- Preventive technique that analyses directly the page fault rate to guess if thrashing is near to happen



- 2nd chance replacement algorithm
- Frame allocation
- Thrashing
- **Frame reservation**

- **Concept**

- Modern OSs keep a percentage of main memory as a store of free frames (**reservation frames**)

- **Goals**

- **To reduce the time taken to serve a page fault**
 - Attempt to have free frames available
 - The replacement algorithm is used:
 - Only when the free frame level gets too low
 - To look for victims to amortize its use
 - Page out
 - Several pages are written at once to disk
- **To avoid thrashing**

- **Reservation frame management**

- The OS guarantees that there will be always a set of free frames
- Some thresholds are set:
 - Minimum number of free frames (M_{MIN})
 - Recommended number of free frames (M_{REC})

$$M_{\text{REC}} \gg M_{\text{MIN}}$$

- Very efficient replacement algorithms are NOT required
 - The first VMS systems used FIFO because their MMU did not have reference bit
 - It is common to use a 2nd chance algorithm (Windows, UNIX SVR4, UNIX 4.4BSD, Linux, HP OpenVMS, ...).

- **Monitor process**

- There is an internal process that periodically **accounts for the number of free frames** (`frame_free`):
 - If `frame_free > Mrec` then do nothing
 - If `frame_free < Mmin` then:
 - Swap out some processes until reaching REC free frames
 - Victims are process that spent more time suspended
 - If `Mmin <= frame_free <= Mrec` then:
 - Seek for processes with too many frames (very low frame rate) to “steal” them some frames applying a replacement algorithm
 - Several victims are selected in every process, the actual number differs on every OS

- **Frame reservation management in thrashing**
 - When thrashing happens the number of free frames decreases quickly
 - The process monitor detects it when:
 - Between two monitor activation **frame_free** decreases too much
 - Solutions:
 - Swap out whole processes until reaching **frame_free** = Mrec
 - OpenVMS and Windows NT
 - Free a constant number of frames in every monitor activation if **frame_free** < Mrec
 - The monitor activation frequency increases
 - UNIX SVR4

- **Reservation frames content**

- The frame of a victim selected by the OS is **not allocated immediately to another page**, instead the frame goes to the **reservation stock**
 - If victim pages are referenced again soon by the process:
 - There is a high probability that they be in the reservation stock and then they can be relocated **without having to read them on disk**
 - The OS remembers which is the content of every frame in the reserve stock
 - If frames included in the reservation stock correspond to **modified pages**:
 - **They are not considered to solve page faults immediately** because its content has to be written on disk (into a file or paging area)
 - » When a **threshold is reached** all these pages are written at once
 - » Page out **overhead is amortized** -> the number of global page writing is minimized
 - After writing on disk, frames become free and can serve page faults

- Virtual memory abstracts physical memory into an extremely large uniform array of storage.
- The benefits of virtual memory include the following:
 - a program can be larger than physical memory,
 - a program does not need to be entirely in memory,
 - processes can share memory,
 - processes can be created more efficiently.
- Demand paging is a technique where by pages are loaded only when they are demanded during program execution.
- A page fault occurs when a page that is currently not in memory is accessed.
- Several algorithms have been analysed: Optima, LRU, approximations to LRU, etc.
- Thrashing occurs when a system spends more time paging than executing.