

# Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informàtica de Sistemes y Computadoras (DISCA)  
*Universitat Politècnica de València*

## Part 2: Process Management

### Seminar 3

## UNIX system calls for processes

fSO

DISCA



UNIVERSITAT  
POLITÈCNICA  
DE VALÈNCIA

- Goals:
  - To know the **services** provided by UNIX **to create processes**
  - To analyse **C language examples** with process related system calls
  - To know the **signal concept**
- Bibliografía:
  - “**UNIX System Programming**”, Kay A. Robbins, Steven Robbins. Prentice Hall. ISBN 0-13-042411-0

- Identification
- Creation
- Waiting
- Ending
- Signals

	Processes
<b>fork</b>	Create a child process
<b>exit</b>	End process
<b>wait</b>	Wait for a process ending
<b>exec</b>	Execute a program
<b>getpid</b>	Get process ID attribute
<b>getppid</b>	Get parent ID attribute

	Signals
<b>kill</b>	Send signals
<b>alarm</b>	Generate an alarm (clock signal)
<b>sigemptyset</b>	Init a mask with no signals set
<b>sigfillset</b>	Init a mask with signals set
<b>sigaddset</b>	Append a signal to a signal set
<b>sigdelset</b>	Delete a signals in a signal set
<b>sigismember</b>	Check if a signals belongs to a signal set
<b>sigprocmask</b>	Check/Modify/Set a signal mask
<b>sigaction</b>	Capture/Manage a signal
<b>sigsuspend</b>	Wait signals capture

- **Identification**
- Creation
- Waiting
- Ending
- Signals

- Every process must have an ID
- The creator process is the *parent* while the created process is the *child*. To know them we use:
  - Process ID (PID) `pid_t getpid(void);`
  - Parent process ID (PPID) `pid_t getppid(void);`

```

/** ej1_getpid.c */
#include <stdio.h>

int main(void)
{
    printf("\nProcess ID: %ld\n", (long)getpid());

    printf("Parent process ID: %ld\n", (long)getppid());

    while(1);

    return 0;
}
    
```

```

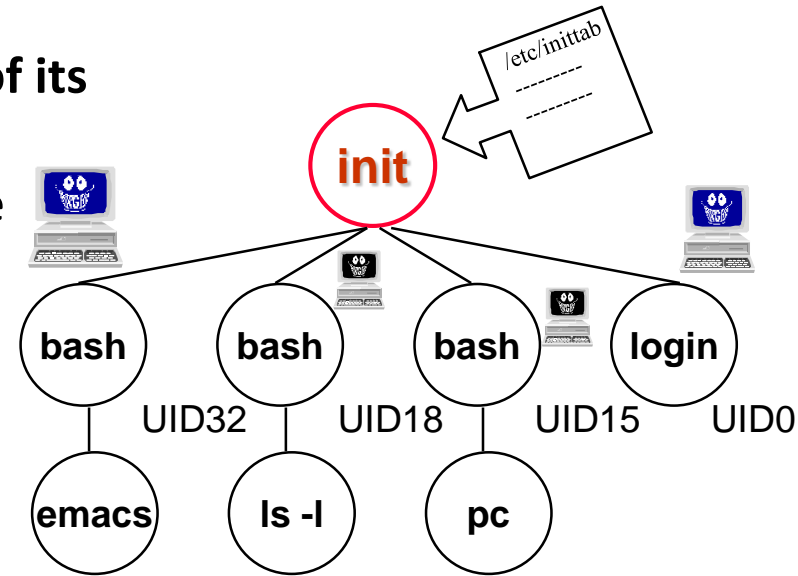
$ gcc -o ej1 ej1_getpid.c
$ ./ej1 &
[1] 2959
$
Process ID: 2959
Parent process ID: 1060
$ ps -l
    
```

UID	PID	PPID	F	CPU	PRI	NI	SZ	RSS	WCHAN	S	ADDR	TTY	TIME	CMD
501	1060	1059	4006	0	31	0	2435548	1088	-	S	ffffff80136e3d50	ttys000	0:00.06	bash
501	2959	1060	4006	0	31	0	2434832	340	-	R	ffffff80140d8300	ttys000	0:04.65	./a.out

- Identification
- **Creation**
- Waiting
- Ending
- Signals

```
pid_t fork(void);
```

- Unix uses a copying mechanism for process creation
  - The **child process is an exact replica of its parent process**
  - The **child process inherits** most of the attributes **from its parent**:
    - Memory image
    - UID, GID
    - Current directory
    - Opened file descriptors
  - Unix assigns an **identifier to every process called PID** at the time of its creation
  - Every process knows the **identifier of its parent process, PPID**
  - The child execution is concurrent with the parent and independent
  - Processes are organized into a **hierarchical tree**





- **fork:** process creation

```
#include <sys/types.h>
#include <unistd.h>

pid_t fork(void)
```

- **Description**

- It creates a child process which is a “clon” from the parent: the child process inherits the most part of the parent attributes.
    - NON Inherited attributes: PID, PPID, pendant signals, file locks and time/accounting.

- **Returning value**

- 0 to the child
    - Child PID to the parent
    - -1 to the parent if an error happens

- **Errors**

- Not enough resources to create the process

Parent and child resume execution with the instruction following `fork()`

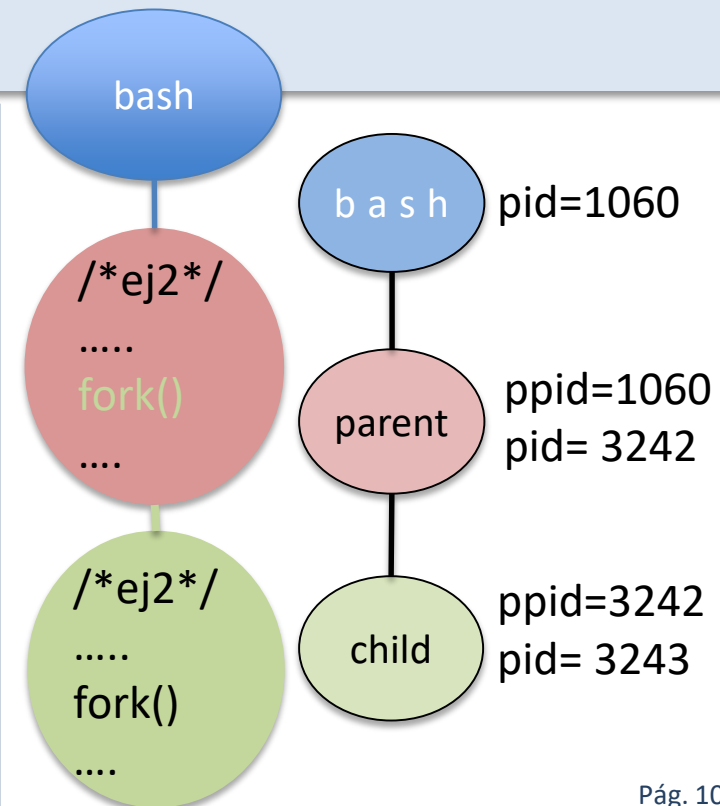
```

/**ej2_fork.c **/
#include <stdio.h>
#include <sys/types.h>

int main(void)
{ printf("Process %ld creates another process\n", (long)getpid());
  fork();
  printf("Process %ld with parent %ld\n", (long)getpid(), (long)getppid());
  sleep(5);
  return 0;
}
    
```

```

$ ps
  PID TTY          TIME CMD
 1060 ttys000    0:00.07  -bash
$ gcc -o ej2 ej2_fork.c
$ ./ej2
Process 3242 creates another process
Process 3242 with parent 1060
Process 3243 with parent 3242
$
    
```



- Parent and child execute different code

```
// ej3_fork.c
#include <sys/types.h>
#include <unistd.h>

int main() {
    pid_t val;
    int var = 0;
    printf("PID before fork(): %d\n", (long) getpid());
    val = fork();
    if (val > 0) {
        printf("Parent PID: %d\n", (long) getpid());
        var++;
    } else {
        printf("Child PID: %d\n", (long) getpid());
    }
    printf("Process [%d]-> var=%d\n", (long) getpid(), var);
    return 0;
}
```

How many processes print this message?

What value(s) of "var" are shown?

- Example

```
// ej4_fork.c

#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>

int main(void)
{
    pid_t pid=fork();

    switch (pid) {
        case -1:
            printf("Could not create child process \n");
            break;
        case 0:
            printf("I am the child with PID %ld and my parent is %ld \n",
                (long)getpid(), (long)getppid());
            break;
        default:
            printf("I am the parent with PID %ld and my child is %ld \n",
                (long)getpid(), pid);
    }
    sleep(5);
    return 0;
}
```

```
$ gcc -o ej3 ej3_fork.c
```

```
$ ./ej3
```

```
I am the parent with PID 3702 and my child is 3704
```

```
I am the child with PID 3704 and my parent is 3702
```

# • Creating processes in chain

```
// ej5_proc_chain.c
#include <stdio.h>
#include <sys/types.h>


#define NPROCESSES 4

int main(void)
{
    pid_t pid;
    int i;

    for (i=0; i<NPROCESSES; i++) {
        pid=fork();
        if (pid!=0)
            break;
        printf("I am the child with PID %ld
               my parent is %ld\n",
               (long)getpid(), (long)getppid());
    }
    sleep(5);
    return 0;
}
```

## Variants

if (pid > 0)  
if (pid==0)  
if (pid<0)



```
$ ./a.out
I am the child with PID 3938 my parent is 3937
I am the child with PID 3939 my parent is 3938
I am the child with PID 3940 my parent is 3939
I am the child with PID 3941 my parent is 3940
```

- **exec()**
  - The call to `fork()` creates a child who is a copy of the calling process
  - What if we want to run a different program? Use **exec()** call
  - There are different versions **exec** depending on the parameters specified:

```
#include <unistd.h>
int execl(const char *path, const char *arg0, ...,
          const char *argn, char * /*NULL*/);
int execlp(const char *file, const char *arg, ... ,
          const char *argn, char * /*NULL*/);
int execl_e(const char *path, const char *arg, ...,
            const char *argn, char * /*NULL*/, char * const envp[]);

int execv(const char *path, char *const argv[] );
int execvp(const char *file, char *const argv[]);
int execve(const char *path, char *const argv[], char *const envp[]);
```

- Variant **l**: arguments are provided separately
- Variant **v**: arguments are provided through a pointer to a vector
- Variant **p**: the location of **file** is searched on the PATH
- Variant **e**: the environment is provided to the child through **envp**, it is not inherited from the parent

- **exec () features**

- It changes the process memory image by the one defined in the new executable file.
- The executable file is specified by its name or its absolute path.
- Some process attributes remain:
  - **Signal handlers, except the ones captured and treated by the default action.**
  - **PID and PPID**
  - **Time accounting**
  - **File descriptors**
  - **Working directory, root directory and file creation mode mask**
- If the SETUID bit of the executable file is set then EXEC sets as effective UID the UID of the executable file owner
  - The same happens with the SETGID bit
- **Errors**
  - Executable file doesn't exist or it is not recognized as executable
  - Permissions
  - Incorrect arguments
  - Not enough memory or resources
- **Returning value**
  - If EXEC returns to the calling program it is because an error happened, the returning value is -1

- Exec() example: the child process lists the current directory

```
// ej6_exec.c
#include <stdio.h>
#include <sys/types.h>

int main(void)
{
    int status;
    pid_t pid=fork();

    char* arguments [] = { "ls", "-l", 0 };

    switch (pid) {
        case -1:
            printf("The child process could not be created \n");
            break;
        case 0:
            printf("I am the child with PID %ld, the current directory content is: \n",
                (long)getpid());
            if (execvp("ls",arguments)==-1){
                printf("Error in exec\n");
                exit(0);
            }
            break;
        default:
            printf("I am the parent process with PID %ld and my child is %d.\n",
                (long)getpid(), pid)
    }
    return 0;
}
```

Another way:

```
execl("/bin/ls", "ls", "-l", NULL)
```



- Identification
- Creation
- **Waiting**
- Ending
- Signals

- The parent has to wait for its child to finish:


```
#include <sys/types.h>
#include <sys/wait.h>
pid_t wait(int *status);
pid_t waitpid(pid_t pid, int *status, int options);
```

- `wait` and `waitpid` stop the calling process execution until a child ends, or until the calling process receives a signal

- `wait()` call

```
pid_t wait(int *status);
```

- It suspends the execution of the calling process until any of its children finishes
  - If there is any zombie child then `wait` returns immediately
- **status:** is not a NULL pointer, it contains information about child termination:
  - Child finished with **exit** or child finished by a **signal**



- **Returning value:**
  - -1 if error or no children
  - Otherwise the finished child PID

- Wait() example:

```
// ej7_wait.c
#include <stdio.h>
#include <sys/types.h>
#include <errno.h>

int main(void)
{
    int status;
    pid_t pid=fork();

    switch (pid) {
        case -1:
            printf(" The child  process could not be created \n");
            break;
        case 0:
            printf("I am the child process with PID %ld and my parent is %ld\n",
                    (long)getpid(), (long)getppid());
            sleep(20);
            printf("I have finished \n");
            break;
        default:
            printf("I am the parent process with PID %ld and my child is %d.
                    Waiting ...\n", (long)getpid(), pid);
            if (wait(&status) != -1)
                printf("My child has ended ok \n");
    }
    return 0;
}
```

- `waitpid()` call

```
pid_t waitpid(pid_t pid, int *status, int options);
```

- It waits for a particular child
- Parameters:
  - **pid**: PID of the child to wait for. If its value is -1 it waits until the first child ends (like **wait**)
  - **status**: like `wait()`
  - **options**:
    - Value 0: blocking (common value)
    - WNOHANG: non blocking
  - **Returning value**:
    - **0** : no child ended (non-blocking version),
    - **-1** : error
    - **> 0** : pid of the returning child


- `waitpid( )` example: fan process creation

```
// ej8_waitpid.c
#include <stdio.h>
#include <sys/types.h>

#define NPROCESSES 4

int main(void)
{
    pid_t pid[NPROCESSES];
    int i, status;

    for (i=0; i<NPROCESSES; i++) {
        pid[i]=fork();
        if (pid[i]==0){
            printf("I am the child %ld my parent is %ld\n",
                (long)getpid(), (long)getppid());
            sleep(10);
            exit(0);
        }
    }
    // Now wait for the third child
    if (waitpid(pid[2],&status,0)==pid[2])
        printf(" My third child has finished \n");
    return 0;
}
```



```
$gcc -o ej8 ej8_waitpid.c
```

- Identification
- Creation
- Waiting
- **Ending**
- Signals

- A process ends completely when:
  - The process itself ends (normally or abnormally) **AND**
  - Its parent process does a wait call
- The normal termination is done by calling to **exit**

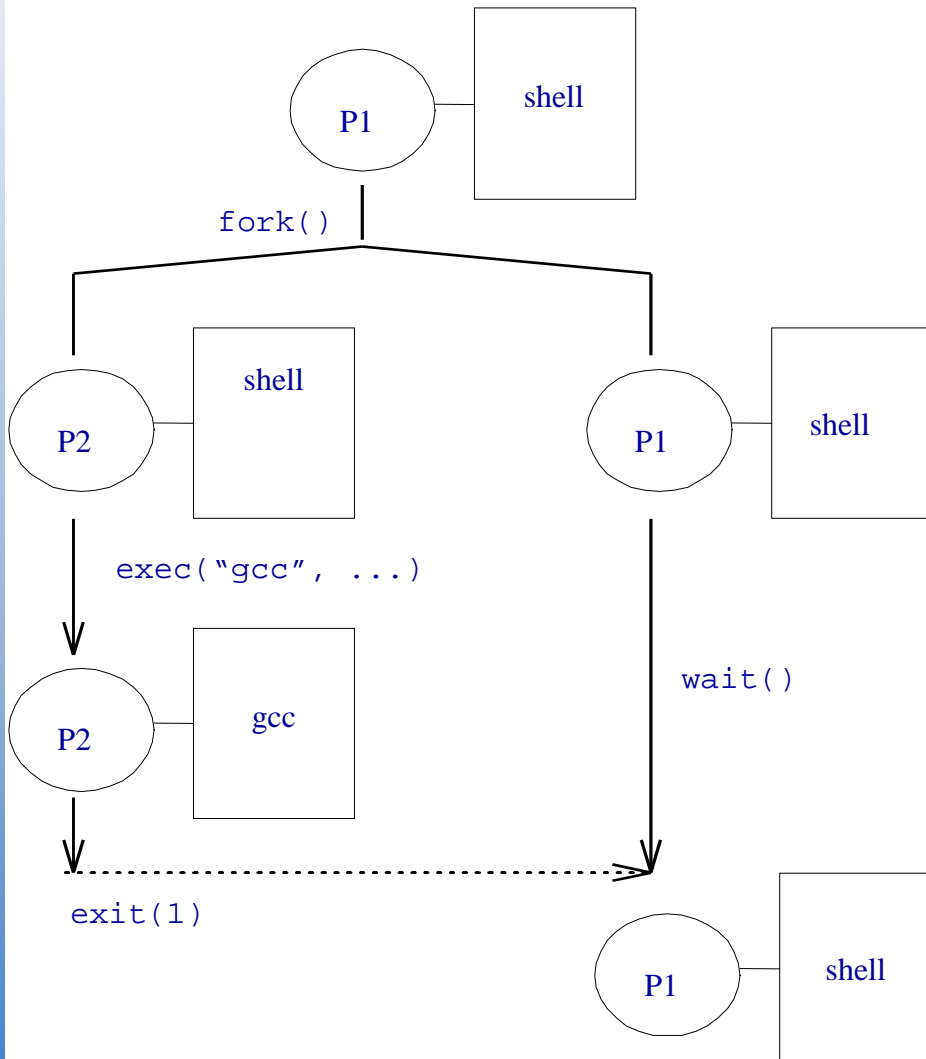
```
void exit (int status)
```

- The status parameter is used to communicate to the parent how the child process ends
- By convention, this value is typically 0 if the process ends normally and any other value means abnormal termination
- The parent process can obtain this value through the system call **wait**



- **Abnormal ending:**
  - The process is ended by the operating system because of an error condition (memory boundary violation, arithmetic errors) or by some other process initiative
    - Signals
- **Zombie:** If the process ends before his parent calls `wait()`:
- **Orphan:** If the parent process terminates before the child:
  - An orphan process is adopted by process root **init**

- Simplified structure of UNIX shell



```


while(TRUE) {
    print_prompt();
    read_command(command, parameters);

    p=fork();           /* creates a child */

    if (p != 0) {       /* parent code */
        waitpid(-1, &status, 0); /* waiting for
                                   the child */
    } else {            /* child code */
        exec(command, params, 0); /* memory image
                                   change */
        error("Command cannot be executed");
        exit(1);
    }
}
  
```

- Identification
- Creation
- Waiting
- Ending
- **Signals**


## Signals generation



- Terminal
- Errors
- Software

`kill(...)`  
`alarm(..)`

## Signal mask



### Mask definition

`sigemptyset(...)`  
`sigfillset(..)`  
`sigaddset(..)`  
`sigdelset(..)`  
`sigismember(...)`

### Mask setting

`sigprocmask(...)`

## Signals handling

`handler(...)`

Handler executes

`SIG_IGN`

Ignors

`SIG_DFL`

Default action (exit)

⋮

### Handler setting

`sigaction(...)`

### Signal waiting

`pause(...)`  
`sigsuspend(...)`

- A signal is the **mechanism used by the OS to inform processes about certain events**
- They all follow the same pattern:
  - It is generated due to the occurrence of an event
  - It is supplied to the process
  - It should be treated by the default handling routine or by the one defined by a specific process
- A signal can:
  - Be handled: a handling routine must be installed
  - Be masked: its handling is delayed
  - Be ignored: its occurrence if not informed