

Fundamentos de los Sistemas Operativos (FSO)

Departamento de Informática de Sistemas y Computadoras (DISCA)

Universitat Politècnica de València

Part 2: Process management

Seminar 6

Synchronization: POSIX Semaphores

fSO

DISCA



UNIVERSITAT
POLITÈCNICA
DE VALÈNCIA


- **Goals:**

- To get use to deal with **critical section** problems
- To know the **synchronization** mechanisms **offered by the OS**
- To use **semaphores** and **mutexes** to solve critical section synchronization

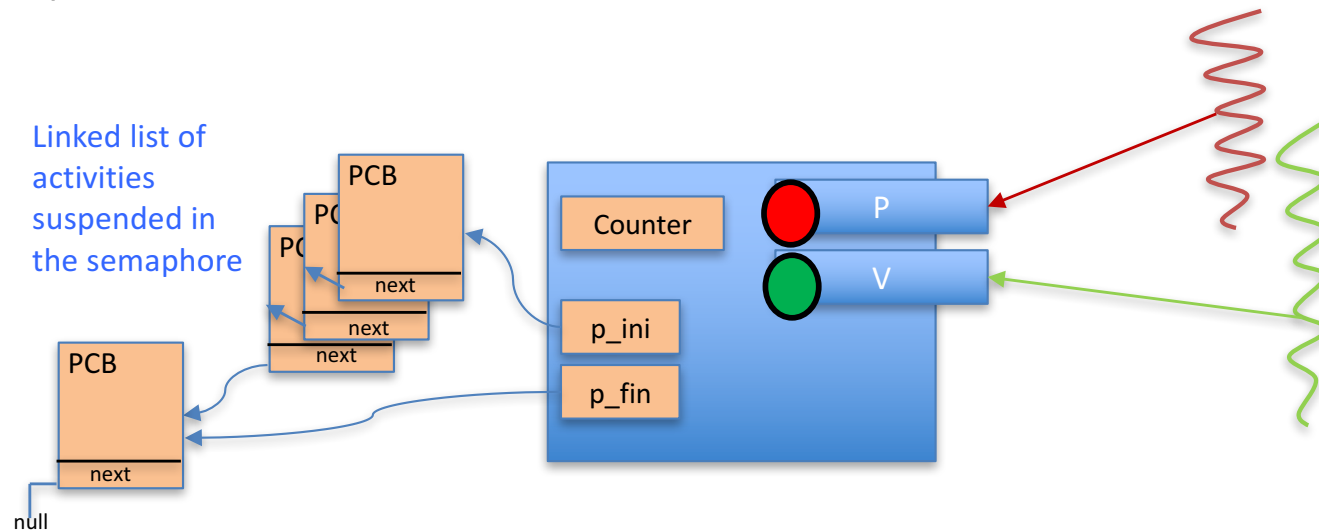
- **Bibliography**

- Silberschatz 8th Ed, chapter 6
- Robbins, chapters 13, 14

- **OS level solutions**
- POSIX semaphores
- POSIX mutexes
- Exercises



- Semaphore
 - Can be seen as an integer that admits an increment by 1 and a decrement by 1 operations performed by an activity (process or thread)
 - The decrement operation can suspend the activity
 - The increment operation can awake another activity previously suspended
 - It is a synchronization object offered by the OS to user activities
 - It is declared as type “semaphore” specifying its initial value
 - The increment (V) and the decrement (P) operations are implemented as system calls



- Initialization, P and V specification

- Declaration and initialization

- Semaphore S(N);**

- It declares a semaphore “S” with “N” as initial value
 - “N” has to be greater or equal to zero

- Decrement

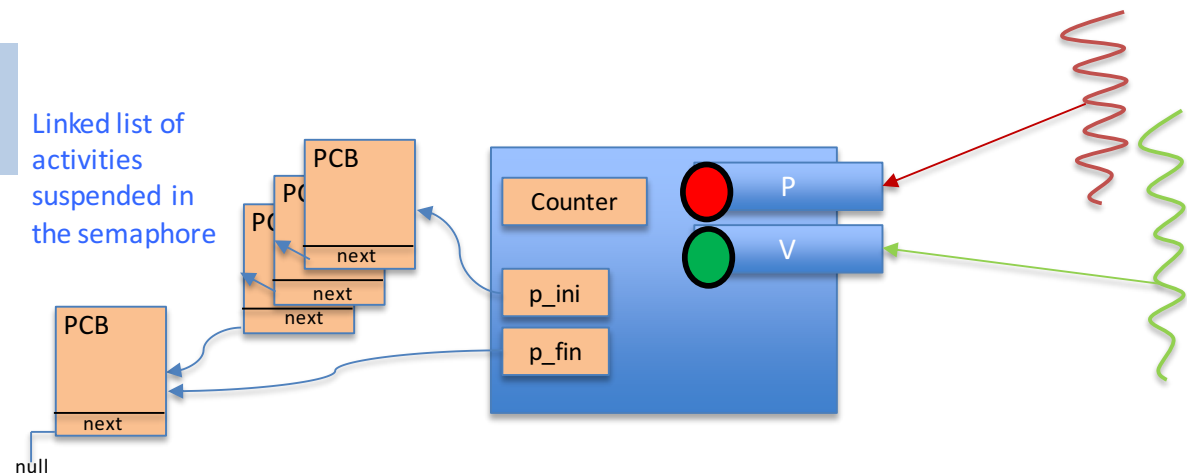
- P(S);**

```
S = S - 1;  
if S < 0 then suspend(S);
```

- Increment

- V(S);**

```
S = S + 1;  
if S <= 0 then awake(S);
```



Notes:

- P and V are guaranteed to be atomic by the OS
- suspend(S) suspends the calling activity in a queue associated to S
- awake(S) extracts an activity from the S queue and awakes it (it goes to the scheduler ready queue)

- Solving the critical section problem
 - We define as a global variable the semaphore “mutex”, initialized to 1
- Semaphore mutex(1);**
- Code for N threads/processes

```
void *thread_i(void *p) {  
  
    while(1) {  
        P(mutex);  
  
        /* Critical section */  
  
        V(mutex);  
  
        /* Remaining section */  
  
    }  
}
```

It complies with:

- Mutual exclusion
- Progress
- Limited waiting, if the semaphore queue is scheduled with FCFS policy

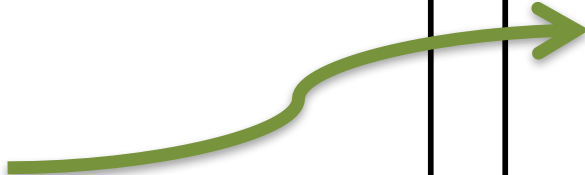
- Another synchronization use of semaphores:
Establishing an execution order
 - We want “thread1” to execute function “F1” before “thread2” executes function “F2”
 - We define a shared semaphore between “thread1” and “thread2” named “sync”, **initialized to 0**

Semaphore sync(0);

```
void *thread1(void *p)
{
    ...

    F1;
    V(sync);
    ...
}
```

```
void *thread2(void *p)
{
    ...
    P(sync);
    F2;
    ...
}
```



- Another synchronization use of semaphores: Limiting the number of activities that can pass through a certain point in the code simultaneously
 - We have a function executed by many threads
 - We want a maximum of 5 threads call function “F” located in a certain place inside the function
 - We define a shared semaphore “max_5”, **initialized to 5**


Semaphore max_5(5);

```
.....  
pthread_t th1, th2, th3, th4, th5;  
pthread_attr_t attr;  
  
pthread_attr_init(&attr);  
pthread_create(&th1, &attr, thread_i NULL);  
pthread_create(&th2, &attr, thread_i NULL);  
pthread_create(&th3, &attr, thread_i NULL);  
pthread_create(&th4, &attr, thread_i NULL);  
.....
```

```
void *thread_i(void *p) {  
  
    ...  
    P(max_5);  
    F;  
    V(max_5);  
    ...  
}
```


- **Basically a semaphore is a “resource counter”**
 - When an activity tries to use a resource, it decrements the counter (**P** operation) and if there are not available resources then the activity is suspended
 - When an activity finishes using a resource then it increments the counter (**V** operation). If there is any suspended activity waiting for the resource, then the **V** operation awakes it.
 - If the counter is greater than 0 then its value indicates the number of available resources
if $S > 0$ then $|S|$ = number of available resources
 - When its value is less than zero its absolute value indicates the number of suspended activities in the semaphore queue
if $S < 0$ then $|S|$ = number of suspended processes
 - When the counter is zero both of the former sentences are true
**if $S = 0$ then \rightarrow there are no suspended processes
there are no available resources**
- With semaphores the OS manages suspended activity queues, waiting for available resources
- Activity synchronization is one of the essential mechanisms that an OS has to provide

- OS level solutions
- **POSIX semaphores**
- POSIX mutexes
- Exercises



- POSIX.1b introduced the semaphore type `sem_t`

```
#include <semaphore.h>
sem_t sem;
```

- A semaphore can be used by all the threads inside a process and also can be shared between processes.
 - After a `fork()` call semaphores in the parent can be inherited by the child depending on “pshared” parameter value in “`sem_init`”.

- POSIX.1b semaphore operations are:

- `int sem_init(sem_t *sem, int pshared, unsigned int value);`
- `int sem_destroy(sem_t *sem);`
- `int sem_wait(sem_t *sem);`
- `int sem_trywait(sem_t *sem);`
- `int sem_post(sem_t *sem);`
- `int sem_getvalue(sem_t *sem, int *sval);`

Operation
P(sem)

Operation
V(sem)

- Producer/Consumer **version 1**

- We define a semaphore “mutex” initialized to 1 →
sem_init(&mutex,0,1);

```
void *func_prod(void *p) {
    int item;

    while(1) {
        item = produce();

        sem_wait(& mutex);

        while (counter == N)
            /*empty loop*/ ;
        buffer[input] = item;
        input = (input + 1) % N;
        counter = counter + 1;

        sem_post(&mutex);
    }
}
```

```
void *func_cons(void *p) {
    int item;

    while(1) {
        sem_wait(&mutex);

        while (counter == 0)
            /*empty loop*/ ;
        item = buffer[output];
        output = (output + 1) % N;
        counter = counter - 1;

        sem_post(&mutex);

        consume(item);
    }
}
```

- If the producer enters the critical section being the buffer full it will get into the while loop forever, the same happens with the consumer when the buffer is empty

- Producer/Consumer **version 2**

```
#include <semaphore.h>
sem_t mutex, items, vacants;
```

```
void *func_prod(void *p) {
    int item;
    while(1) {
        item = produce();
        sem_wait(&vacants);
        sem_wait(&mutex);

        buffer[input] = item;
        input = (input + 1) % N;
        counter = counter + 1;

        sem_post(&mutex);
        sem_post(&items);
    }
}
```

```
void *func_cons(void *p) {
    int item;
    while(1) {
        sem_wait(&items);
        sem_wait(&mutex);

        item = buffer[output];
        output = (output + 1) % N;
        counter = counter - 1;

        sem_post(&mutex);
        sem_post(&vacants);
        consume(item);
    }
}
```

```
sem_init(&mutex, 0, 1);
sem_init(&vacants, 0, N); // indicates the initial number of buffer vacants
sem_init(&items, 0, 0); // indicates the initial number of buffer items
...
```

- OS level solutions
- POSIX semaphores
- **POSIX mutexes**
- Exercises



- **Mutex**

- POSIX.1c defines the **mutex** object for thread synchronization
 - It can be seen as **semaphores that can only be initialized to 1**
 - It is used only to guaranty mutual exclusion
 - It works as a latch -> two operations: **lock** and **unlock**
 - A mutex has at every moment:
 - **State**: Open or closed
 - **Owner**: It is a thread that has executed a successful **lock** operation on it



- Mutex operation:
 - A mutex is created opened and without owner
 - When a thread calls to **lock**
 - If the mutex was open (without owner), the calling thread closes it and becomes its owner
 - If the mutex was closed, the calling thread is suspended
 - When an owner thread calls to **unlock**
 - The mutex gets open
 - If there were suspended threads on it, one of them is awaked, then the mutex gets closed and the awoken thread is the new owner

- **Mutex operation**

A mutex is created **opened** and without owner



When a thread **calls to lock** and



when the **owner** thread calls to **unlock** and ...



mutex is open?

NO



YES



the calling thread **closes it** and becomes its owner



the calling thread is **suspended**



there are suspended threads?

NO

YES



one of them is awaked, then the mutex gets closed an the awoken thread is the **new owner**



The mutex gets **opened**



- **POSIX calls for mutexes**
 - Creation and destruction
 - pthread_mutex_init
 - pthread_mutex_destroy
 - Attribute initialization
 - pthread_mutexattr_init
 - pthread_mutexattr_destroy
 - Changing/Checking attribute values
 - Lock and unlock
 - pthread_mutex_lock
 - pthread_mutex_trylock
 - pthread_mutex_unlock

Equivalent
to P(sem)

Equivalent
to V(sem)

- **Example:** concurrent access to a shared variable by two threads
 - Main function code:

```
pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;

int V = 100;

// Threads code (next slide)

int main ( ) {
    pthread_t      thread1, thread2;
    pthread_attr_t  attributes;
    pthread_attr_init(&attributes);
    pthread_create(&thread1, &attributes, func_thread1, NULL);
    pthread_create(&thread2, &attributes, func_thread2, NULL);


    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
}
```

- Example (cont)
 - Thread functions:

```
void *func_thread1(void *p) {  
    int c;  
  
    for(c=0; c<1000; c++) {  
        pthread_mutex_lock(&m);  
  
        V = V + 1;  
  
        pthread_mutex_unlock(&m);  
    }  
    pthread_exit(0);  
}
```

```
void *func_thread2(void *p) {  
    int c;  
  
    for(c=0; c<1000; c++) {  
        pthread_mutex_lock(&m);  
  
        V = V - 1;  
  
        pthread_mutex_unlock(&m);  
    }  
    pthread_exit(0);  
}
```

- OS level solutions
- POSIX semaphores
- POSIX mutexes
- **Exercises**



- **Exercise S06.1** What possible values will take **x** as a result of the concurrent execution of the following threads?

```
#include <semaphore.h>
#include <pthread.h>
#include <stdlib.h>
#include <stdio.h>
sem_t s1,s2,s3;
int x;
```



```
void *func_thread1(void *a)
{
    sem_wait(&s1);
    sem_wait(&s2);
    x=x+1;
    sem_post(&s3);
    sem_post(&s1);
    sem_post(&s2);
}
void *func_thread2(void *b)
{
    sem_wait(&s2);
    sem_wait(&s1);
    sem_wait(&s3);
    x=10*x;
    sem_post(&s2);
    sem_post(&s1);
}
```

```
int main()
{
    pthread_t h1,h2 ;
    x = 1;
    sem_init(&s1,0,1); /*Inicializa a 1*/
    sem_init(&s2,0,1); /*Inicializa a 1*/
    sem_init(&s3,0,0); /*Inicializa a 0*/

    pthread_create(&h1,NULL,func_thread1,NULL);
    pthread_create(&h2,NULL,func_thread2,NULL);
    pthread_join(h1,NULL);
    pthread_join(h2,NULL);
}
```

Exercise S06.2 What possible values will take shared variables **x** and **y** at the end of the following concurrent threads. The initial values are: $x=1$, $y=4$, $S1=1$, $S2=0$ y $S3=1$.

Thread A

```
P(S2);  
P(S3);  
x = y * 2;  
y = y + 1;  
V(S3);
```

Thread B

```
P(S1);  
P(S3);  
x = x + 1;  
y = 8 + x;  
V(S2);  
V(S3);
```

Thread C

```
P(S1);  
P(S3);  
x = y + 2;  
y = x * 4;  
V(S3);  
V(S1);
```

