



Unit 3:

Database Management Systems (DBMS)

Unit 3. Database Management Systems

1. The ANSI/SPARC Architecture

- 1.1. Schemas
- 1.2. DBMS fundamentals
- 1.3. Data independence

2. Transactions, Integrity, and Concurrency

- 2.1. Transactions
- 2.2. Semantic integrity
- 2.3. Concurrent access control

3. Recovery and Security

- 3.1. DB Recovery
- 3.2. Security

1.1. Schemas

Original proposal

Proposal of a DBMS architecture by the working group ANSI/SPARC (1977). They proposed the database definition with 3 levels of abstraction:

- **Internal level** → Internal schema
Description of the DB in terms of its physical representation
- **Conceptual level** → Conceptual schema
Description of the DB independently of the DBMS. It is usually a graphical representation.
- **External level** → External schemas
Description of the different users' partial views

Refined proposal

Since there was no generalized conceptual model for different kinds of DBMS (it is difficult to obtain the physical data structures from a conceptual graphical representation) the “logical” level was added:

- **Internal (physical) level** → Internal schema (*not in this course*)
DB description in terms of its physical representation. Describe how the database is store in secondary memory.
- **Logical level** → Logical schema (Unit 1)
DB description in terms of the DBMS data model. It does not include any details of the physical representation.
- **Conceptual level** → Conceptual schema (Unit 4)
Description of the information system from the organizational point of view. It is independent of the DBMS. We will use a UML class diagram.
- **External level** → External schemas (authorizations and views)
Description of the partial views which the different users have on the DB.

Example: Logical schema

Department (cod_dep: char(4), nombre: char(50), teléfono: char(8), director: char(9))

PK:{cod_dep}

NNV:{nombre}

FK:{director} -> Lecturer(dni) On delete set nulls. On update cascade

Subject (cod_asg: char(5), nombre: char(50), semestre: char(2), cod_dep: char(4),
teoría: real, prácticas: real)

PK:{cod_asg}

NNV:{nombre, semestre, cod_dep, teoría, prácticas}

Uni:{nombre}

FK:{cod_dep} -> Department(cod_dep)

On delete restrict. On update cascade

IC₁:(teoría <= prácticas)

IC₂:(semestre IN {'1A','1B','2A','2B','3A','3B','4A','4B'})

Lecturer (dni: char(9), nombre: char(80), teléfono: char(8), cod_dep: char(4),
provincia: char(25), edad: entero)

PK :{dni}

NNV :{nombre, cod_dep}

FK :{cod_dep} -> Department(cod_dep)

On delete restrict. On update cascade

Teaching (dni: char(9), cod_asg: char(5), gteo: entero, gpra: entero)

PK :{dni,cod_asg}

NNV :{gteo,gpra}

FK :{dni} -> Lecturer(dni) On delete cascade. On update cascade

FK :{cod_asg} -> Subject(cod_asg) On delete restrict. On update cascade

General constraint: **GC1**: "All teacher must lecture at least one subject".

Example: Internal schema

Depends on the DBMS.

Subject:

Hash file by cod_dep

B+ index over (semestre + cod_dep)

Lecturer:

Hash file by nombre

Department:

Hash file by cod_dep

B+ index over nombre

Teaching:

Disordered file

Example: External schema

External schema for the **Maths Department**

```
CREATE VIEW Maths-Lecturer AS
    SELECT dni, nombre, teléfono, categoría, edad
    FROM Lecturer
    WHERE dptocod_dep = 'DMA';

CREATE VIEW Maths-Subject AS
    SELECT cod_asg, nombre, semestre, teoría, prácticas
    FROM Subject
    WHERE cod_dep = 'DMA';

CREATE VIEW Maths-Teaching AS
    SELECT T.dni, T.cod_asg, T.gteo, T.gpra
    FROM Lecturer L, Teaching T, Subject S
    WHERE L.cod_dep = 'DMA'
        AND S.cod_dep = 'DMA'
        AND L.dni = T.dni
        AND S.cod_asg = T.cod_asg;
```

A DBMS that supports the 3-level architecture must:

- Allow the **definition** of the different **schemas** for the database (except the conceptual schema),
- Establish the **correspondence** between schemas,
- **Isolate the schemas**: changes in one schema should not affect neither the schemas at upper levels nor the application programs.

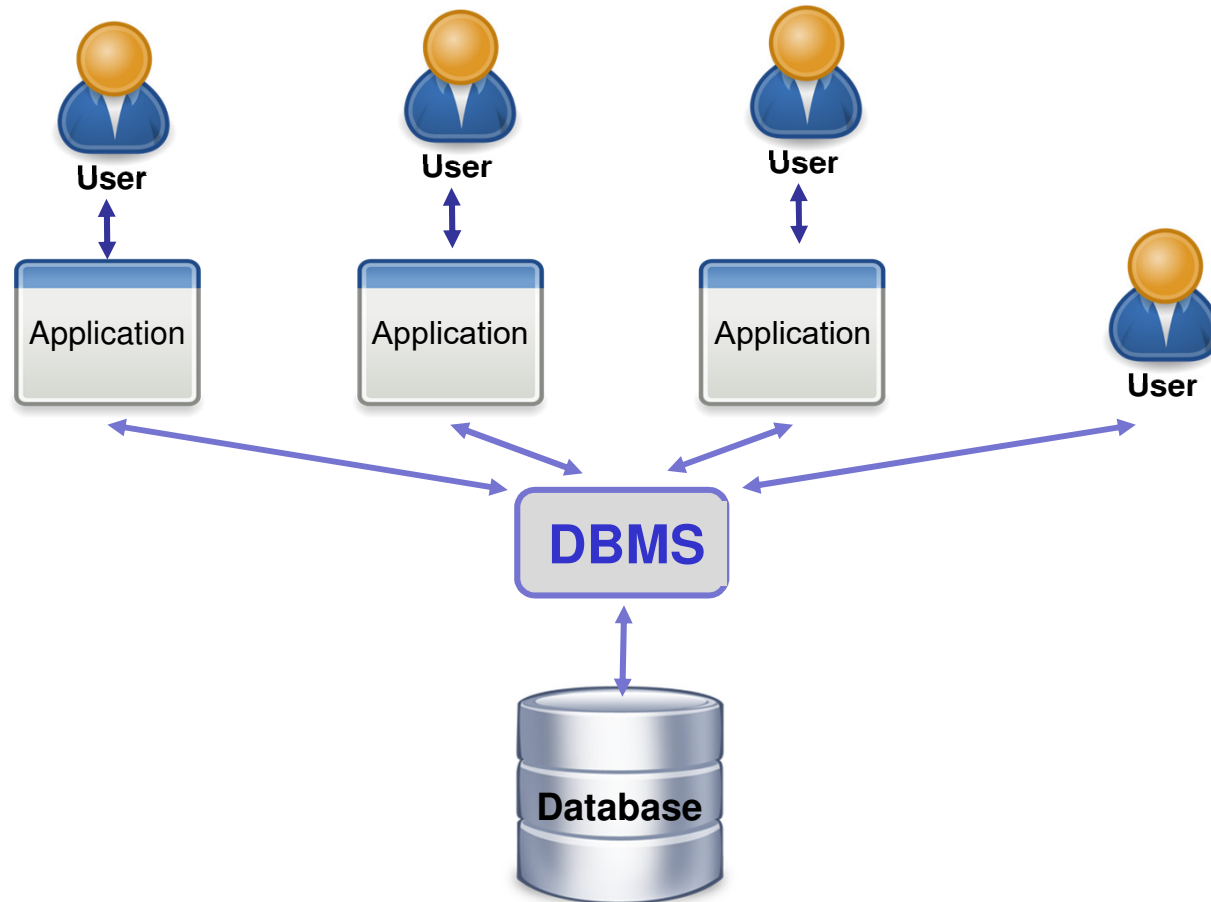


DATA INDEPENDENCE

1.2. DBMS Fundamentals

Functions of a DBMS

DBMS: Software which allows the creation and manipulation of databases (DB).



A DBMS must maintain the independence, integrity and security of data.

Functions of a DBMS

Objectives of DB techniques

- Unified and independent **data description**
- **Application** independence
- Partial **view** definition

DBMS Functions

Data definition at several **levels**

- Logical schema
- Internal schema
- External schema

DMBS Components

Schema definition languages and their associated translators

Functions of a DBMS

Objectives of DB techniques

- Information **management**

DBMS Functions

Data manipulation

- Query
- Update

Management and **administration** of the database

DMBS Components

Manipulation languages and their associated translators

Tools for:

- Restructuring
- Simulation
- Statistics
- Printing and reporting

Functions of a DBMS

Objectives of DB techniques

- Data **integrity** and **security**

DBMS Functions

Control of:

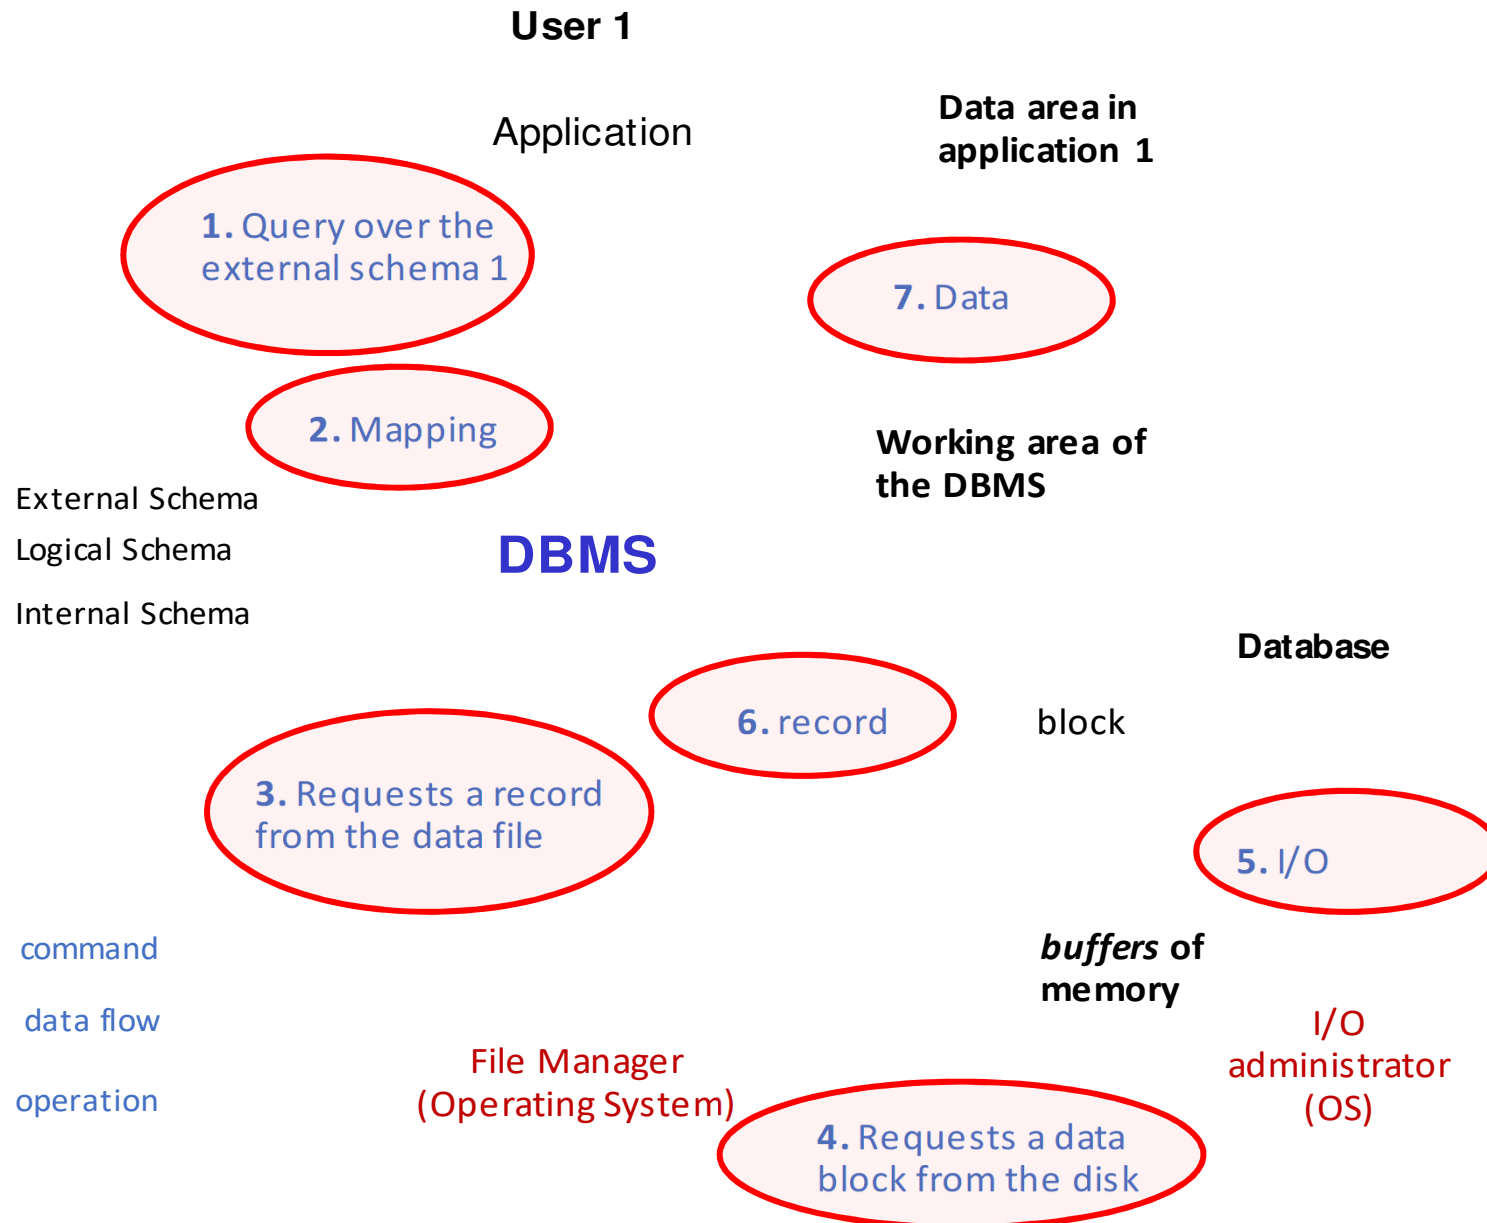
- Semantic integrity
- Concurrent access
- Recovery in case of failure
- Security (privacy)

DMBS Components

Tools for:

- Integrity control
- Reconstruction
- Security control

Accessing the data



Binding

Binding (*ligadura*):

Transformation of the external schema into the internal schema.

Types

- Logical binding (steps 2 and 7).
- Physical binding (steps 3 and 6).

When the binding is performed, independence disappears



It is important to determine the binding moment

Currently, most DBMS do the binding for each query.

1.3. Data Independence

Data independence

Property which ensures that the application programs are independent of

- the changes which are performed on the logical schema corresponding to data which they do not use

or

- the physical representation details of the accessed data

Advantages of Data independence

- Helps to improve the quality of the data.
- Database system maintenance becomes affordable.
- Enforcement of standards and improvement in database security.
- It is not necessary to alter data structure in application programs.
- Developers can focus on the general structure of the Database rather than worrying about the internal implementation.
- Easily make modifications in the physical level is needed to improve the performance of the system.

Logical independence

Logical independence between the logical schema and the external schemas:

The external schemas and the application programs cannot be affected by the modifications in the logical schema of data which are not used by these programs

EXAMPLE:

If we add new attributes to the “*Department*” table, such as the date in which the department was created, the building,... the external schema of the “*Maths-department*” does not need to be modified.

Physical independence

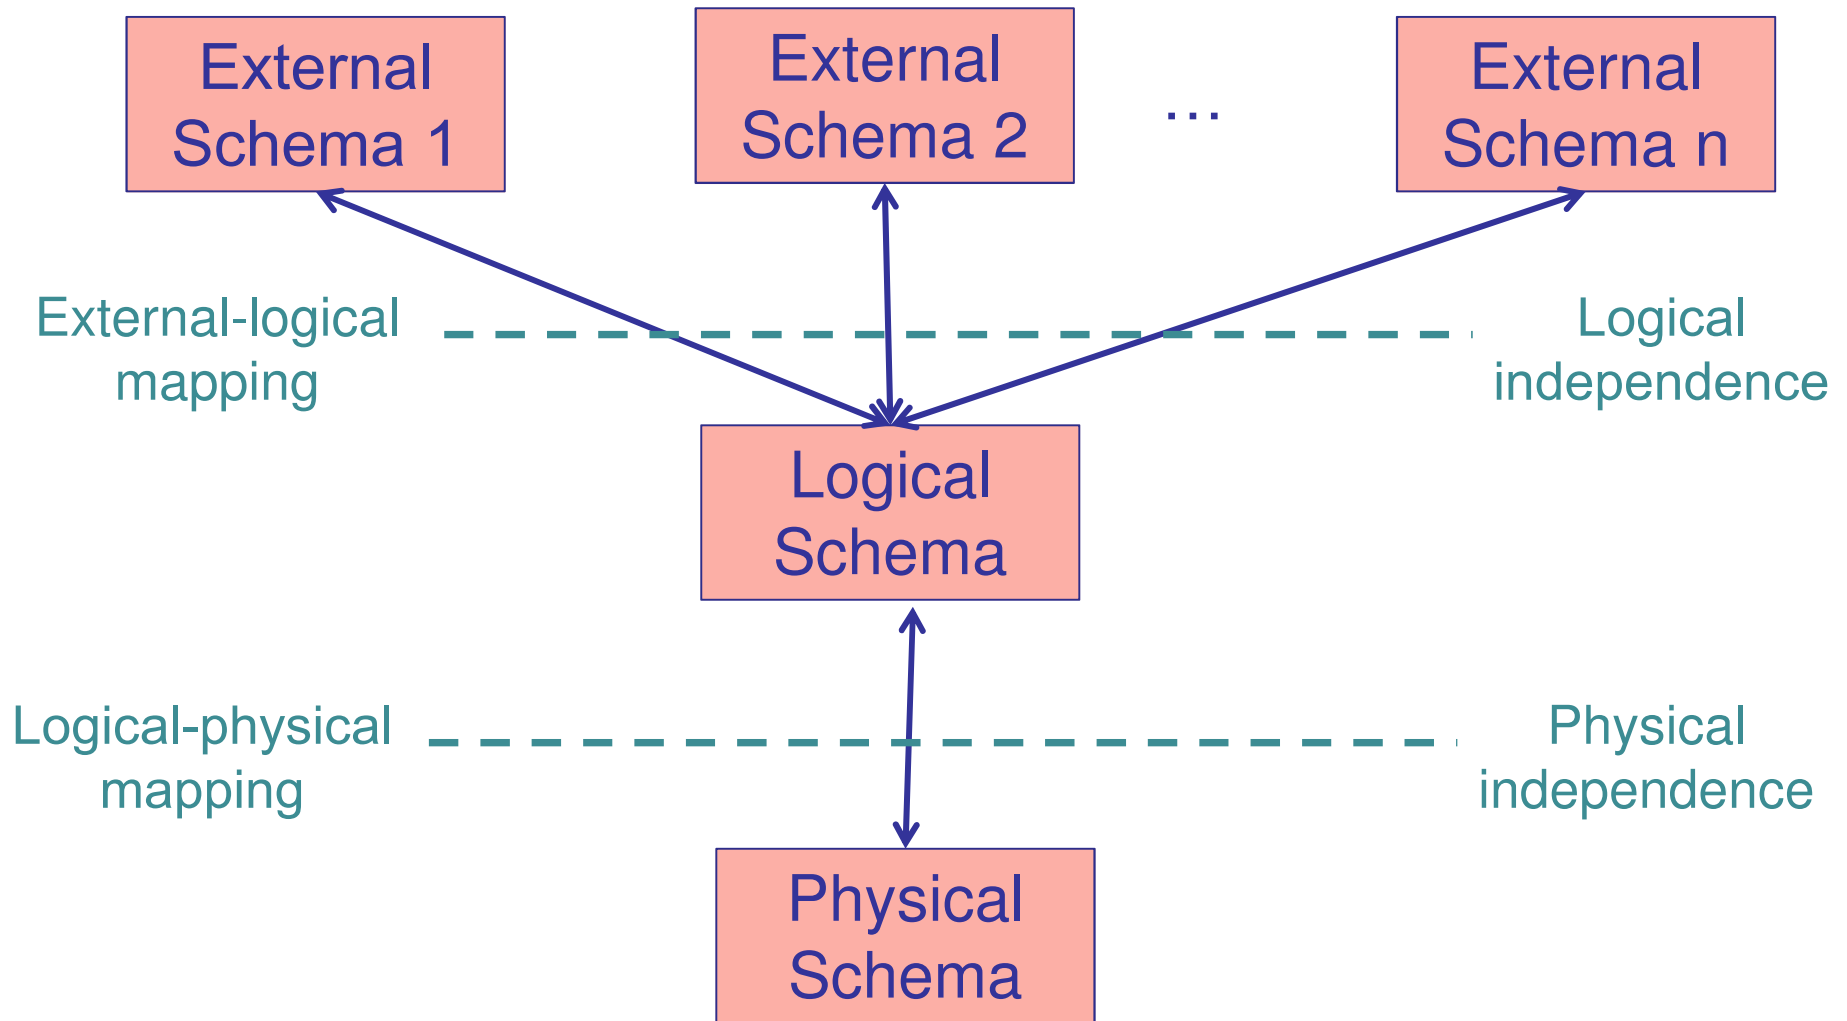
Physical independence between the internal schema and the logical schema:

The logical schema cannot be affected by changes in the internal schema which refer to the implementation of the data structures, access modes, page size, search path, etc.

EXAMPLE:

If the data structures used in the implementation of the “*Subject*” table are changed, the logical schema does not need to be modified.

Data independence



Physical independence is found in most DBMS, while logical independence is more difficult to find.

Unit 3. Database Management Systems

1. The ANSI/SPARC Architecture

1.1. Schemas

1.2. DBMS fundamentals

1.3. Data independence

2. Transactions, Integrity, and Concurrency

2.1. Transactions

2.2. Semantic integrity

2.3. Concurrent access control

3. Recovery and Security

3.1. DB Recovery

3.2. Security

2. Transactions, Integrity, and Concurrency

Objective of DB technology



Information quality

*“Data must be structured in such a way to adequately **reflect** the **objects**, **relations**, and **constraints** which exist in the part of the real world modeled by the database model.”*

When reality changes → User **updates** the database

The information contained in the DB must preserve the schema definition.

2. Transactions, Integrity, and Concurrency

Information quality (integrity perspective) means that:

- The DBMS must ensure that the data are **correctly stored**
- The DBMS must ensure that **user updates** over the DB are correctly executed and become **permanent**.

2. Transactions, Integrity, and Concurrency

DBMS Tools oriented towards integrity:

- Check (when an update is performed) the **integrity constraints** defined in the schema.
- Control the correct execution of the **updates** in a **concurrent** environment.
- Recover (**reconstruct**) the DB in case of loss or accident

2.1. Transactions

2.1. Transactions

The operations in a DB are organized in transactions.

Transaction:

Sequence of access operations to the DB which constitute a **logical execution unit**.

Example

Emp (id, name, address, dept)

PK: {id}

FK: {dept} → Dep(code)

Dep (code, name, location)

PK: {code}

IC₁: All departments have at least one employee

Insert a new department:

<“d2”, “Human Resources”, “2nd floor”>

whose first employee is the id 20

Example

1st
Idea

1) Insert in *Dep*:

<d2, "Human Resources", "2nd floor">

ERROR: IC₁ is violated

2) Modification of *Emp* on the tuple with *id* 20

2nd
Idea

1) Modification of *Emp* on the tuple with *id* 20

ERROR: the FK over dept in Emp is violated

2) Insertion in *Dep*:

<d2, "Human Resources", "2nd floor">

Defining transactions

Actions which change transactions states:

Begin:

Indicates the **beginning** of the execution of the transaction

Cancellation (user rollback):

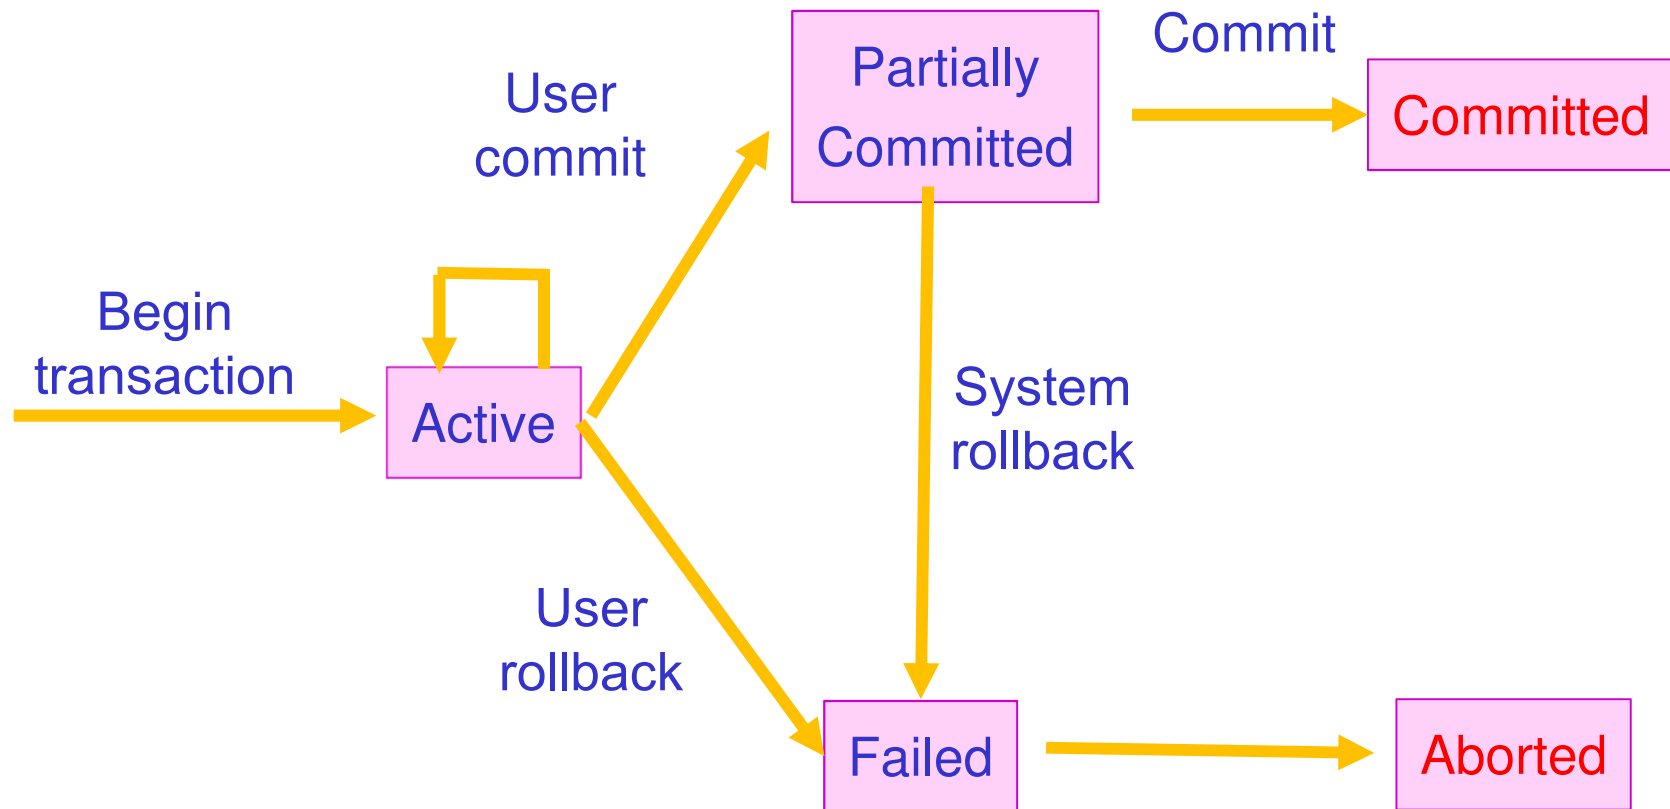
The user aborts the transaction.

Confirmation (*user commit*):

The user considers the transaction as ended. Then the DBMS performs some checking to determine how the transaction will end:

- **Success (system commit):**
Indicates the **success** of the transaction, making the DBMS store the changes performed on the DB.
- **Failure (system rollback):**
Indicates the **failure** of the transaction, or that the transaction hasn't passed the checking. The DBMS undoes all the possible changes performed by the transaction.

States of a transaction



Properties of Transactions (ACID)

Atomicity:

A transaction is an indivisible unit that is either **performed** in its **entirety** or is not performed at all (“All or nothing”).

Consistency:

The transaction must transform the DB from one consistent state **to another consistent state** (all integrity constraints must be met)

Isolation:

Concurrent transactions execute independently: All the partial effects of **incomplete transactions should not be visible** to other transactions

Durability:

The effects of a **successfully** completed (committed) transaction are **permanently recorded** in the DB and must not be lost because of a subsequent system or other transaction failure

2.2. Semantic Integrity

2.2. Semantic Integrity

Integrity constraint:

Property of the real world which is modelled by the DB

- Constraints are defined in the **logical schema** and the DBMS must ensure that they are met.
- Checking is **performed** whenever the **DB changes** (when any updating operation is executed)
- Constraints **not included in the DB** schema must be maintained by the application programs

*This situation is, in general, **inappropriate** if the constraints are common to more than one application, since the responsibility to check them is dispersed*

Types of integrity constraints

Static: They must be met in each state of the DB (they can be represented by logical expressions)

Example:

The *lab credits* in a *subject* cannot be greater than the *lectures credits*.

Dynamic (Transition): They must be met regarding two consecutive states.

They are not usually implemented by commercial DBMS

Example:

The *credits* of a *subject* cannot decrease.

Expressing static constraints in SQL

- Constraints over possible **data values**. e.g. Domains
- Constraints over **attributes**. e.g. NNV.
- Constraints over **relations**. e.g. PK, FK
- **General constraints** over the DB. e.g: “*All subject must be lectured by at least one teacher.*”

When are checked:

- After every command (**IMMEDIATE**)
- At the end of the transaction(**DEFERRED**)

Expressing transition constraints in SQL

Triggers (*“disparadores”*)

- Using triggers, the designer can program the system response when some events are produced.
- This allows to incorporate complex constraints into the DB.

A trigger includes:

1. **Events**: Operations over the DB which trigger it
2. **Conditions** to determine if the actions must be executed or not.
3. **Actions** to be executed when an event happen and the conditions are met. They are usually written in a data-oriented high level programming language, that can include SQL commands.

Expressing transition constraints in SQL

Example:

A trigger to implement the integrity constraint: *“A lecturer can only teach a subject assigned to his/her department”*

1. **Events:** INSERTion of a tuple in the “Subject” table
2. **Conditions:** The lecturer and the subject are not in the same department.
3. **Actions:** Reject the operation (the insertion).

2.3. Concurrent access control

2.3. Concurrent access control

In order to keep the integrity of the database, the DBMS must control concurrent access to the database:

To avoid that the results of the execution of **several processes** (users or programs) lead to **incorrect**, **incoherent** or **lost results** because of the simultaneous execution of other program accessing the same data

Basic operations

Basic operations in a transaction which are relevant to the DBMS:

read(X):

Reading or access to a piece of data X in the DB over the program variable with the same name.

write(X):

Update (insertion, deletion, or modification) of a piece of data X in the DB by using the program variable with the same name

Reading steps

read(X):

1. Seek the address of the block which contains the datum X
2. Copy the block to a buffer into main memory
3. Copy the datum X from the buffer to the program variable X

Writing steps

write(X):

- If not read before {
1. Find the **address of the block** containing the datum X
 2. **Copy** the block into a database **buffer** in main memory
 3. **Copy** the **datum** X from the program variable into the database **buffer**
 4. **Write the updated block** from the database buffer to the disk

Possible problems

The DBMS must control the concurrent access by the applications.

Problems due to interference of concurrent accesses:

- a) **Loss of updates.** An apparently successfully completed update operation by one user can be overridden by another user.
- b) **Inconsistent** information corresponding to several valid database states. One transaction reads several values but a second transaction updates some of them during the execution of the first.
- c) Access to updated data (but **still not confirmed**) that can still be cancelled. One transaction is allowed to use the intermediate results of another transaction before it has committed (*dirty read*).

A. The lost update problem

Time	Program 1	Program 2	
t1	read(11548, teoría) teoría=4,5		<i>Two programs reading and updating the theory credits of subject 11548:</i> <ul style="list-style-type: none"> • P1 adds 1,5 credits • P2 adds 2 credits
t2		read(11548, teoría) teoría=4,5	
t3	teoría←teoría+1,5 teoría=6		
t4		teoría←teoría+2 teoría=6,5	
t5	write(11548, teoría)		
t6		write(11548, teoría)	

Subject

cod_asg	nombre	semestre	cod_dep	teoría	prácticas
11545	Análisis Matemático	1A	DMA	4,5	1,5
11546	Álgebra	1B	DMA	4,5	1,5
11547	Matemática Discreta	1A	DMA	4,5	1,5
11548	Bases de Datos y Sistemas de Información	3A	DSIC	6,5	1,5

← 4,5 + 1,5 + 2 = 8

B. The inconsistent analysis problem

Program 1: List for each lecturer his/her amount of credits

Teaching				Time	Program 1	Program 2
dni	cod_asg	gteo	gpra	t1	Calculate credits for lecturer 111 Credits 111= 9 (1×4,5 + 3×1,5)	Change a theory group of subject 11545 from lecturer 564 to lecturer 111
111	11545	2	3	t2	Calculate credits for lecturer 123 Credits 123= 9 (0×4,5 + 2×1,5 + 1×4,5 + 1×1,5)	
123	11545	0	2	t3		
123	11547	1	1	t4	Calculate credits for lecturer 453 Credits 453= 0	
564	11545	1	2	t5	Calculate credits for lecturer 564 Credits 564= 7,5 (1×4,5 + 2×1,5)	
Subject						
cod_asg	...	teoría	prácticas			
11545	...	4,5	1,5			
11546	...	4,5	1,5			
11547	...	4,5	1,5			
11548	...	4,5	1,5			

Result:	
DNI	Credits
111	9 credits
123	9 credits
453	0 credits
564	7,5 credits

C. The uncommitted dependency problem

Subject

cod_asg	...	teoría	prácticas
11545	...	4,5	1,5
11546	...	4,5	1,5
11547	...	4,5	1,5
11548	...	6	1,5

Dirty read

Time	Program 1	Program 2
t1	read(11548, teoría)	
t2	teoría ← teoría + 1,5	
t3	write(11548, teoría)	
t4		read(11548, teoría)
		teoría = 6
		<i>Use this value (teoría = 6) in its instructions</i>
t5		
t6		confirmation
t7	cancellation	

Techniques

Reserving some data occurrences (locks)

- In examples **a)** and **c)** P1 must lock a record.
- In examples **b)** all the table must be locked.
- Need for controlling deadlocks

Other solutions (for the example c): Cascade cancellation or transaction isolation

Unit 3. Database Management Systems

1. The ANSI/SPARC Architecture

1.1. Schemas

1.2. DBMS fundamentals

1.3. Data independence

2. Transactions, Integrity, and Concurrency

2.1. Transactions

2.2. Semantic integrity

2.3. Concurrent access control

3. Recovery and Security

3.1. DB Recovery

3.2. Security

Recovery and Security

A database must guarantee:

- **Recovery:**
*(Part of integrity, but not from the point of view of consistency.
Recovery is focus on durability and persistence)*

A database must always be recovered from any type of failure.

- **Security:**
A database cannot allow non-authorized access.

3.1. DB Recovery


DB recovery

The transaction properties of **atomicity** and **durability** force a DBMS to ensure that:

- If **confirmed**, the changes performed are recorded in the DB to make them persistent.
- If **cancelled**, the changes performed over the DB are **undone**.

Database recovery: Example

Backups are not sufficient



4-10-2021: Update on accounts

- Transaction #1: Update
- Transaction #2: Update
- Transaction #3: Update
- ...
- Transaction #51: **System failure !**

Recovery Procedure:

- Replace the file “Accounts” with its backup

Negative effect:

- The updates of 50 transactions are lost

DB Recovery

Backups alone are **not** the **solution** to the recovery problem.

- The increase of the backup frequency is not a feasible solution.

DB technology provides much more efficient and robust techniques for DB recovery.

Lost of confirmed data is inadmissible
with current technology.

Causes of transaction failure

1. **Local to the transaction** (normal system operation)

- **Transaction error** (incorrect DB access, failed calculations, etc.)
- **Exceptions** (integrity violation, security problems, etc.)
- **Concurrency control** (locked state between two transactions)
- **Human decision** (inside a program or explicitly).

2. **Extern to the transaction** (system error)

- A. System failures with **loss of main memory**.
- B. Failures in the storage system with **loss of secondary memory** (disk failure, human errors, virus infection,...)

A. Failures of system main memory

- The changes performed by a transaction are located in memory buffers (main memory).
- When the transaction is **confirmed** its changes must be recorded in secondary memory.
- If a failure with loss of main memory occurs between the **transaction confirmation** and **flushing the buffers** to secondary memory, the blocks in the buffers will be lost.

B. Failures of secondary memory

- The changes performed by a **confirmed** transaction are recorded into the DB
- If there is a **failure in secondary memory**, the changes will be lost.

3.1.1 Recovery from failures of system main memory

Recovery from failures of system main memory

What to do:

- Recover confirmed transactions which have not been recorded.
- **Cancel** transactions which have **failed**.

Who:

Recovery module

How:

Most used technique: Use a *journal file* (or *log*, “*fichero diario*”).

Transaction implementation

Two ways of implementing transactions:

- Transactions with Immediate Update
Updates **can** have an **immediate effect on secondary memory**. In case of cancellation, they have to be undone.
- Transactions with Deferred Update
Updates will only have immediate effect on main memory.
The updates will be transferred to secondary memory **when confirmed**.
- Both need a log file

Log (journal) file

Activities and events are recorded in the log file:

- **Record** the update operations performed by existing transactions during a period of time (e.g. one day)
- The log (journal) file is stored on **disk** to avoid loss after a system failure.
- It is **dumped periodically** into a massive storage unit (magnetic tape, optical disk,...).

Types of entries in a log file

[**start**, T]: A transaction with identifier T has been started.

[**write**, T, X, value_before, value_after]: The T transaction has performed an update instruction on data X.

[**read**, T, X]: The T transaction has read data X.

[**confirm**, T]: The T transaction has been confirmed.

[**cancel**, T]: The T transaction has been cancelled.

Problems:

- Size of log file can increase very quickly.
- Recovery in case of failure is very expensive (many instructions have to be redone).

Solution:

checkpoints (*“puntos de control o verificación”*)

Checkpoints → They are recorded in the log file periodically

How it works ?

- **Suspend** the execution of transactions temporally.
- **Record** a checkpoint in the log file.
- **Record** all updates performed by confirmed transactions (copy all main memory **buffers to disk**).
- **Resume** the execution of the suspended transactions..

A. IMMEDIATE DATABASE UPDATES

Updates **can** have an **immediate effect on secondary memory**.
In case of cancellation, they have to be undone

Failure of a transaction $T \rightarrow$ **Undo** changes performed by T

Update the data which has been modified by T with its original value (*value_before*): Search for the entries in the logfile [**write**, T , X , *value_before*, *value_after*]

Unconfirmed transactions

[start, T] in the log file without *[confirm, T]*

→ Undo changes performed by *T* (previous process)

Confirmed transactions

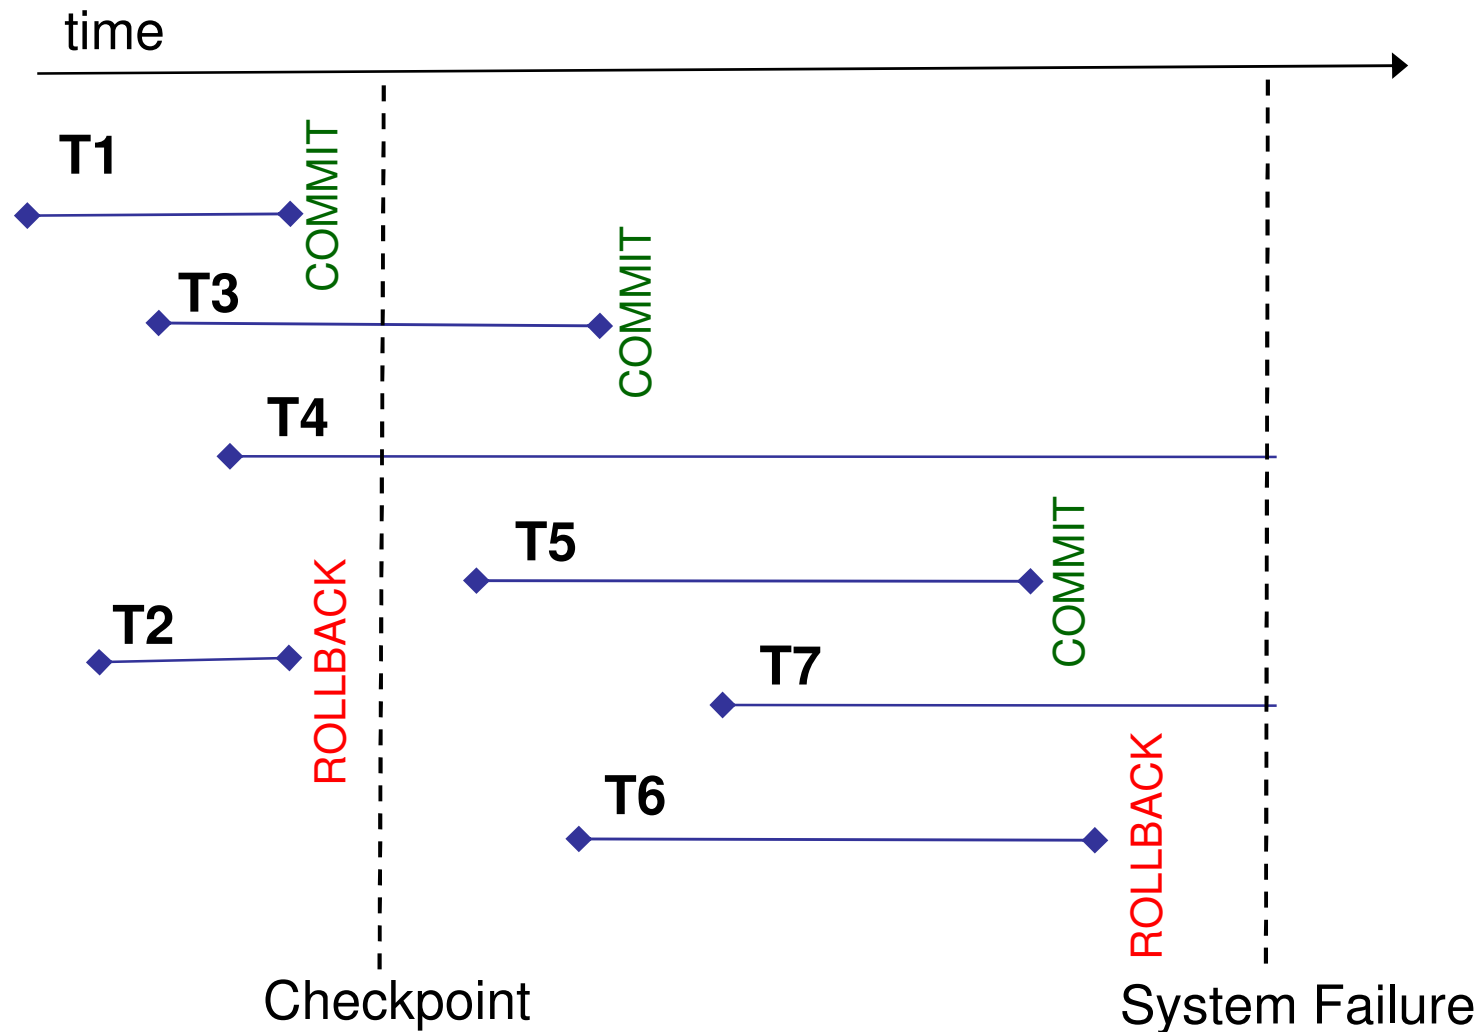
[confirm, T]

→ Execute (redo) them again:

[write, T, X, value_before, value_after]

DB recovery with **immediate** Updates

Recovery the DB from the last checkpoint



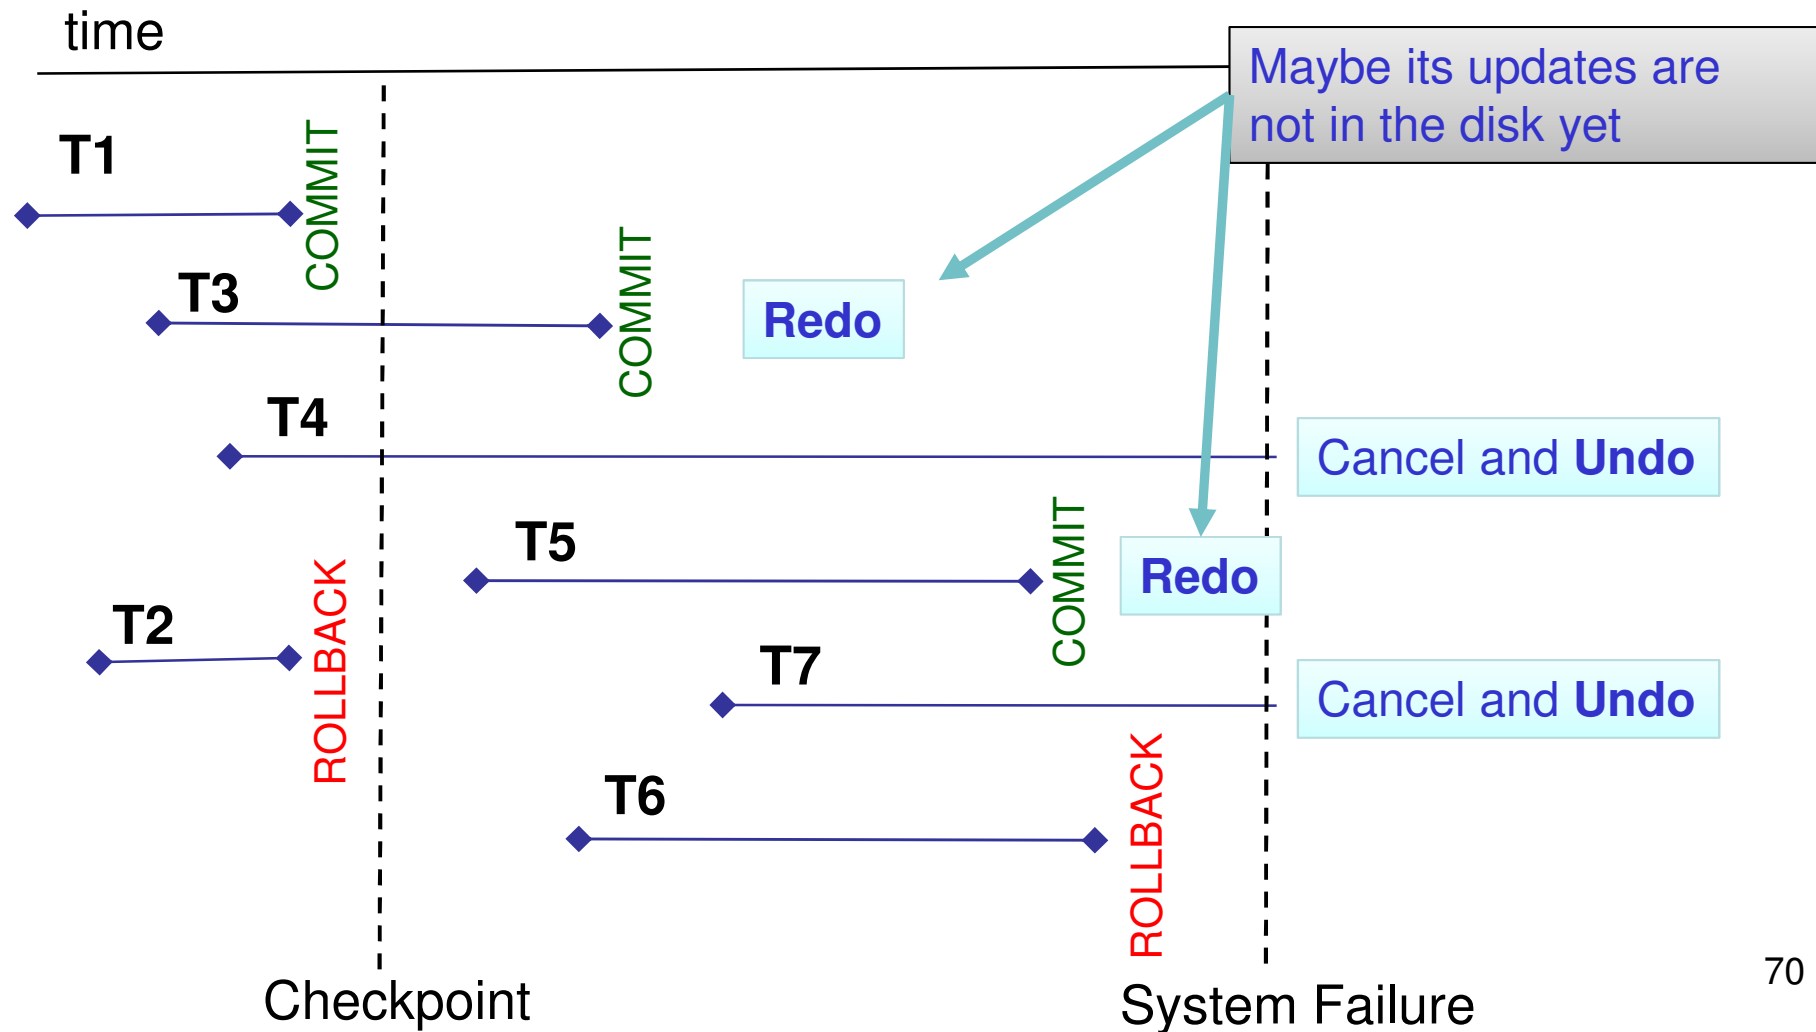
DB recovery with **immediate** Update

Basic considerations:

- **Updates** performed by **confirmed** transaction (*a transaction commit appears in the log file*) could have **not** been transferred to **disk** because the buffer block where they are, has not been recorded yet: **Redo**
- **Updates** performed by **non-confirmed** transactions (there is no transaction commit in the log file) could be **in disk** because their main memory blocks were transferred to disk: **Undo** (only used in immediate updates)
- When a **checkpoint** is recorded, the DBMS **records** all the updates performed by the **confirmed** transaction.

DB recovery with **immediate** Update

Recovery the DB from the last checkpoint



B. DEFERRED DATABASE UPDATES

Updates only have immediate effect on main memory. The updates will only be written to the secondary memory **when confirmed** (after the commit).

- **Unconfirmed** transactions

[start, T] in the log file without *[confirm, T]*

→ **do nothing** they are not in secondary memory

- **Confirmed** transactions

[confirm, T]

→ **Execute (redo)** them again:

[write, T, X, value_before, value_after]

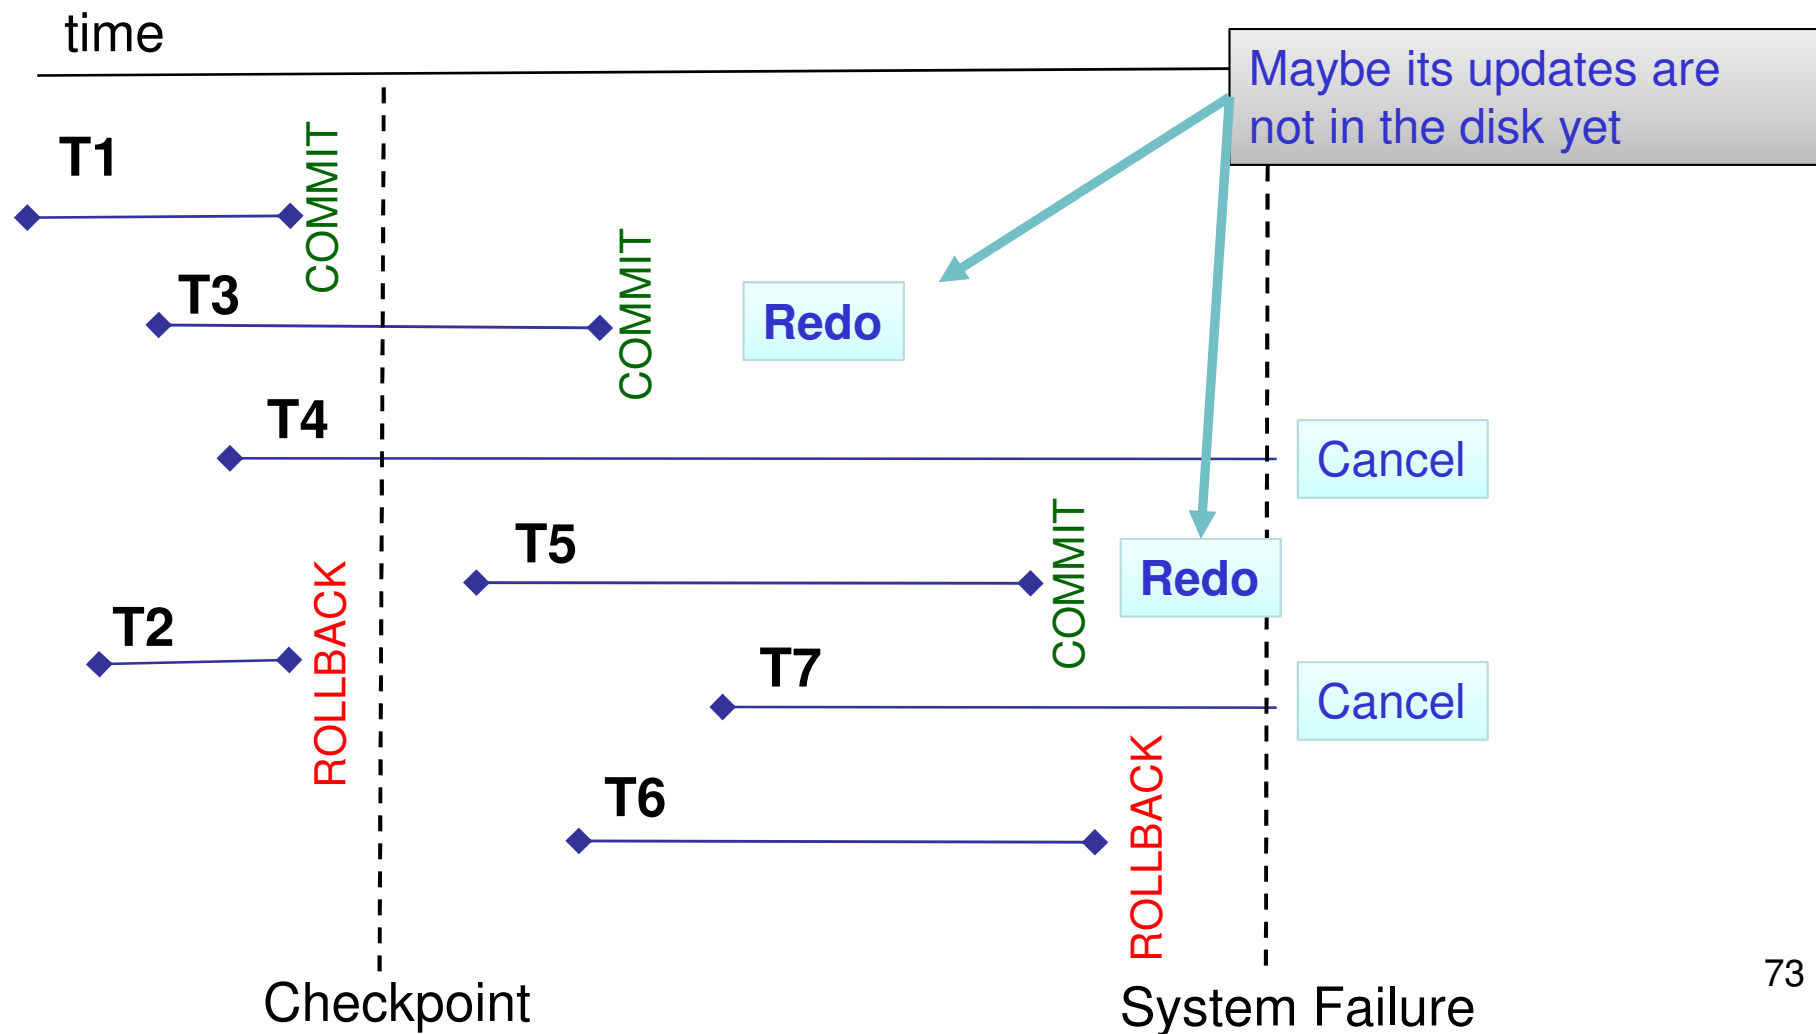
DB recovery with deferred Update

Basic considerations:

- Updates performed by confirmed transaction (*a transaction commit appears in the log file*) could have not been transferred to disk because the buffer block where they are, has not been recorded yet: Redo
- Updates performed by non-confirmed transactions (there is no transaction commit in the log file) are not in disk : **Do nothing**
- When a checkpoint is recorded, the DBMS records all the updates performed by the confirmed transaction.

DB recovery with **deferred** Update

Recovery the DB from the last checkpoint



3.1.2 Recovery from failures of secondary memory

Recovery from failures of secondary memory

When a failure of the **storage system** occurs, the database might be **damaged** totally or partially.

Technique:

Reconstruction of the database:

- Using the most recent **backup**
- From the backup instant, the system uses the **log file** to redo all the instructions performed by the confirmed transactions.

3.2. DB Security

3.2. Security

Objective:

Information can only be accessed by the people and processes that are authorized and in the authorized way

Techniques

- User identification
- Establishment of allowed accesses:


Modes:

- **Authorization list** associated to each user containing the allowed objects and operations. (*GRANT*)
 - **Level of authorization** (less flexible). There several users groups with different authorization.
-
- Management of transferrable authorizations:
Handover of authorizations from one user to another. (*WITH GRANT OPTION*)

Management of transferrable authorizations


It is necessary to know the access **authorizations** of each **user** (some authorizations will be transferable to other users).

- One authorization can be transferred to other user in **mode transferable or not**




When one authorization is **revoked**:

- If the authorization was transferable, it is necessary to revoke all the transferred authorizations



if one user receives more than one authorization, each of the authorizations can be independently revoked.



3.2.1. Privacy and Security

Ethical and legal implications

Extreme care on:

- Protection against the access or spreading of **personal data** to **non-authorized users**.
- Control the **flow** to third parties of **information** that can contain personal data or **information** that is apparently **aggregated** (parameterized queries) but that might **reveal particular** information for some parameters.
- **Custody of security backups**, retired or malfunctioning disks, etc.
- **Small devices** (USB disks or sticks, smart phones, tablets, etc.): **lost** or **stolen** very easily.

Personal Data:

- Personal data is “any information relating to an identified or identifiable natural person.”
- In the European Union: The EU General Data Protection Regulation (GDPR) became enforceable across Europe on 25 May 2018
- Spain “***Agencia de Protección de Datos***”.

Exercises

1.- The physical schema of a database is modified due to a change of disks. If an application uses a view that in turn uses tables that are stored in any of these disks, what happens to the application?

Justify the answer according to the ANSI/SPARC architecture and the concept of independence.

2.- Consider two transactions T1 and T2, which are running concurrently, both working on a piece of data X. Indicate whether one or more properties of transactions are not satisfied here. Briefly justify the answer.

	T1	X in T1	T2	X in T2
Time ↓	read(X)	5		
	X=X+1	6		
	write(X)	6		
		6	read(X)	6
	confirm	6	...	6
	...			

3. How can be recovered a DB from a main memory failure using a logfile and checkpoints ? (Assume immediate updates)

4. Consider the following time diagram, and assuming a DBMS with immediate update. If a system failure occurs, as illustrated in the following figure

- What should the DBMS do if the system failure is a main memory loss?
- What should the DBMS do if the failure is a secondary memory loss ?

